

國立東華大學應用數學系

碩士論文

*Exploration and Extension in Spatial Extremes Incorporating Thin-plate
Splines Basis Functions*



研究生：謝易軒
指導教授：黃灝勻 博士

中華民國 115 年 6 月

國立東華大學應用數學系

碩士論文

待定

To Be Determined



研究生：謝易軒

指導教授：黃灝勻 博士

中華民國 115 年 6 月

學位考試委員會審定書

國立東華大學 _____ 系所

研究生 _____ 君所提之論文

經本委員會審查並舉行口試，認為

符合碩士學位標準。

學位考試委員會召集人 _____ 簽章

指導教授 _____ 簽章

委 員 _____ 簽章

委 員 _____ 簽章

委 員 _____ 簽章

系主任 _____ 簽章
(所長)

中華民國 _____ 年 _____ 月 _____ 日

國立東華大學
NATIONAL DONG HWA UNIVERSITY

學位論文原創性聲明書
DECLARATION OF THESIS/DISSERTATION ORIGINALITY

學位論文題目：

Thesis/Dissertation Title：

本人在此聲明，所呈交的學位論文是在指導教授黃灝勻的指導下，由個人獨立研究所完成之最終版本。本人對論文內容負責，除了文中已經標註引用處的內容外，論文不包含任何其他他人已經發表或撰寫過的研究成果。對本研究及學位論文做出重要貢獻的個人和組織，均已在文中以明確方式標明。

該論文內容如有違反學術道德或學術規範的行為，如造假、變造、抄襲、研究成果重複發表或未適當引註、以違法或不當手段影響論文審查、不當作者列名等，本人願意承擔由此而產生的法律責任和法律後果。

I declare that the thesis/dissertation herein is the final version of my work, which is composed and accomplished individually under the guidance of my supervisor, Prof. Hao-Yun Huang. I am responsible for the contents of this thesis/dissertation: It contains no research result that was previously published or written by another person. Information derived from published and unpublished work of others has been acknowledged in the text, and a list of references is given. Any contribution made by other individual or organization is explicitly acknowledged in the thesis/dissertation.

If any research misconduct, including fabrication, falsification, or plagiarism in proposing, performing, or reviewing research, or in reporting research results, is discovered in my thesis/dissertation, I am willing to bear corresponding legal responsibilities and all the results therefrom.

聲明人 Declarant：

日期 Date：_____ (yyyy/mm/dd)

1140303修訂

摘要

我的論文的目的延續 USA EPA，是利用 AQS 校正 CMAQ 的預測值。CMAQ 的尾部預測較差。Morris 用 skew-t process 及 matern covariance modeling gaussian noise term 來校正 CMAQ 的預測值。Morris 還在建模上設計有 time-varying coefficients with AR(1) structure, censoring to the tail, partition 的開關讓使用者更有彈性去見模。我的論文的目的在於將 time-varying coefficients expand to AR(2) structure, and incorporate multiple resolution thin-plate splines basis functions as covariates to improve modeling non-stationarity spatial dependence. 給予使用者更多的彈性去見模，並且提升模型的預測能力。

關鍵字：



Abstract

To more accurately calibrate CMAQ by the ozone site data, Morris develop partition, threshold, AR(1) time-varying coefficients based on spatial skew-t process. To let model more scalable, we extend the time-varying coefficients to AR(2) structure. And incorporate the multiple thin-plate spline basis functions as covariates to capture the non-stationary spatial dependence. We find that in some level of exceedance, the model with the AR(2) time varying coefficients or the mrts covariates have a lower brier score than the Morris model. Give user more choice in model selection.

Keywords:



Contents

學位考試委員會審定書	i
學位論文原創性聲明書	ii
摘要	iii
Abstract	iv
Contents	v
1 Introduction	1
2 Methodology	1
2.1 Spatio-temporal skew-t process	1
2.2 MRTS covariates	1
2.3 Extension to AR(2) time-varying coefficients	1
3 Simulation study	1
3.1 Design	1
3.2 Cross validation	2
3.3 Results	3
4 Ozone data analysis	4
4.1 Model comparisons	5
4.2 Results	6
5 MERRA-2 data analysis	8
5.1 Data preprocessing and model setup	8
5.2 Results	9
6 Discussion	10

1 Introduction

2 Methodology

2.1 Spatio-temporal skew- t process

2.2 MRTS covariates

2.3 Extension to AR(2) time-varying coefficients

3 Simulation study

3.1 Design

The simulation study is used to evaluate whether the proposed extensions improve predictive performance beyond the space-time skew- t threshold model of Morris et al. (2017). The benchmark simulation follows the structure of Morris et al.: observations are generated at $n_s = 144$ spatial sites, sampled uniformly on $[0, 10] \times [0, 10]$, over $n_t = 50$ independent time points. The mean structure is constant in the data generation, $\mathbf{X}^\top \boldsymbol{\beta} = 10$, while the fitted models include an intercept and first-order spatial trend. The main data-generating settings include Gaussian data, skew- t data with one knot, skew- t data with five knots, a max-stable setting, and a Brown-Resnick max-stable setting. For each setting, 50 simulated data sets are generated.

For the baseline comparison, five Bayesian models are considered:

1. Gaussian marginal distribution with $K = 1$ knot;
2. Skew- t marginal distribution with $K = 1$ knot and no thresholding;
3. Symmetric- t marginal distribution with $K = 1$ knot and threshold $T = q(0.80)$;
4. Skew- t marginal distribution with $K = 5$ knots and no thresholding;
5. Symmetric- t marginal distribution with $K = 5$ knots and threshold $T = q(0.80)$.

Here $q(0.80)$ denotes the empirical 80th percentile of the simulated response. These five models represent the main combinations of tail behavior, thresholding, and spatial partitioning used in the Morris model.

Two additional simulation experiments are conducted for the extensions in this thesis. First, to evaluate the AR(2) time-varying coefficient structure, data are generated with temporal dependence in the latent τ , $\mathbf{w}$, and $\mathbf{z}$ processes. The experiment uses $n = 100$ sites and $T = 30$ time points, with 90 sites used for training and 10 sites withheld for prediction. We repeat the experiment 20 times and compare three temporal structures: no time series, AR(1), and AR(2). Each fitted model uses a skew- t marginal distribution with $K = 5$ knots and no thresholding, and the MCMC is run for 5000 iterations with 2500 burn-in iterations.

Second, to evaluate the MRTS covariate extension, we augment the original five models with multi-resolution thin-plate spline basis functions (Tzeng and Huang, 2018). These basis functions are added to the regression component in order to capture nonstationary spatial structure that may not be fully represented by a stationary Matérn covariance. We focus on simulation settings 4 and 5: setting 4 is a one-knot skew- t data-generating process with $\lambda = 3$, and setting 5 is a five-knot partitioned skew- t data-generating process with $\lambda = 3$. For setting 4, each original model is compared with versions using 10, 15, and 20 MRTS covariates. For setting 5, each original model is compared with a version using 15 MRTS covariates. Each setting uses 50 simulated data sets, and each MCMC chain is run for 20000 iterations with 10000 burn-in iterations.

3.2 Cross validation

Predictive performance is evaluated by withholding spatial sites from model fitting and scoring posterior predictive distributions at the held-out sites. In the benchmark simulation, 100 sites are used as training sites and 44 sites are withheld for validation. In the AR(2) experiment, the split is 90 training sites and 10 testing sites. The MRTS experiment follows the benchmark simulation structure and compares each MRTS-augmented model with the corresponding non-MRTS model on the same simulated data sets.

Because the main inferential target is prediction of high-level exceedances, the primary scoring rule is the Brier score (Gneiting and Raftery, 2007). For an exceedance level L , the Brier score is $BS(L) = \{I(y > L) - \hat{P}(Y > L)\}^2$, where y is the held-out observation and $\hat{P}(Y > L)$ is the posterior predictive probability of exceeding L . The score is averaged over all held-out sites and time points, and lower values indicate better exceedance prediction.

For the extension experiments, we also record MAPE and quantile scores. The quantile score at level α is used to evaluate the accuracy of posterior predictive quantiles, with lower values again indicating better prediction. The reported exceedance and quantile levels are concentrated in the upper tail: the AR(2) experiment uses $q(0.90)$, $q(0.95)$, $q(0.98)$, and

$q(0.99)$, while the MRTS experiment uses levels from $q(0.90)$ to $q(0.995)$.

3.3 Results

The benchmark simulation supports the main conclusion of Morris et al. (2017): when the fitted model matches the data-generating mechanism, it tends to improve predictive performance, and when the data are non-Gaussian the skew- t models generally improve over a Gaussian model. Thresholded models are not uniformly better in the simulation setting, which is expected because most simulated data are generated directly from the full skew- t model rather than from a censored tail model. The simulation therefore confirms that the skew- t framework provides flexibility for asymmetric and heavy-tailed data while remaining competitive with Gaussian methods when the data are light-tailed and symmetric.

To integrate the proposed extensions in one comparison framework, we summarize AR(2) and MRTS evidence together. For AR(2), Table 1 reports selected quantitative results from the detailed missing-at-location experiments in `ar2_extend_detail_doc.md`. Across scenarios, AR(2) sometimes improves mid-tail Brier scores (for example BS90/BS95 in stronger-dependence settings), but AR(1) or no-time-series can still be preferable for runtime or at specific tail levels. This reinforces that temporal flexibility does not imply uniform predictive dominance.

Scenario (missing-at-location)	BS90 (AR2/AR1/NoTS)	BS95 (AR2/AR1/NoTS)	BS99 (AR2/AR1/NoTS)	Time(s) (AR2/AR1/NoTS)
σ^2 AR(2) only	0.0129/0.0130/0.0128	0.0106/0.0106/0.0106	0.0032/0.0032/0.0032	54.78/54.53/51.41
All large ϕ	0.0344/0.0346/0.0352	0.0269/0.0267/0.0268	0.0071/0.0061/0.0069	138.47/135.74/109.64
All small ϕ	0.0267/0.0267/0.0266	0.0178/0.0178/0.0179	0.0059/0.0060/0.0060	232.74/228.68/187.20
Big ϕ_τ , others small	0.0388/0.0390/0.0389	0.0181/0.0182/0.0182	0.0045/0.0044/0.0043	135.17/134.59/102.93

Table 1: Selected AR(2) simulation results under missing-at-location settings. Values are moved from `ar2_extend_detail_doc.md`; lower values indicate better performance.

For MRTS, the simulation evidence is more pattern-based than scenario-wise dominant: in setting 4 and setting 5, adding MRTS covariates can improve some upper-tail Brier or quantile scores, but the gains are not stable across all model families and quantile levels. Increasing the number of MRTS covariates also increases computation time, especially when $K_{\text{MRTS}} = 20$. Table 2 places AR(2) and MRTS side-by-side as proposed extensions in a common decision summary.

Overall, the simulation study indicates that both extensions expand the usable model class,

Extension	Simulation setup	Main empirical pattern	Practical implication
AR(2) time-varying coefficients	Missing-at-location experiments with AR(2), AR(1), and no-time-series fitted to multiple dependence regimes	AR(2) can improve selected mid-tail Brier scores, but runtime advantages are often with AR(1) or no-time-series; no uniform winner across regimes	Use AR(2) as a candidate temporal model, not as a default replacement
MRTS covariates	Benchmark simulation settings 4–5 with additional multi-resolution thin-plate spline covariates ($K_{\text{MRTS}} = 10, 15, 20$ or 15)	Improvements appear at selected tail levels, but average differences are small and not stable across quantiles/model classes; runtime increases with larger K_{MRTS}	Use MRTS selectively when nonstationarity is plausible and tail-score gains justify added cost

Table 2: Integrated simulation summary for the two proposed extensions.

but neither extension uniformly dominates the Morris baseline. The AR(2) structure provides temporal flexibility, and MRTS covariates enrich spatial nonstationarity representation; however, their empirical benefits depend on data-generating conditions and target exceedance levels. Therefore, in the ozone analysis, AR(2) and MRTS should be treated as candidate options to be selected by held-out predictive scores.

Recommended additional experiments to strengthen logical completeness:

(1) Add a full AR(2) vs AR(1) vs NoTS experiment under the exact benchmark simulation design ($n_s = 144$, $n_t = 50$, same train/test split as Morris), so temporal-structure comparisons are directly comparable with the baseline five-model benchmark.

(2) Add a factorial AR(2) $\times$ MRTS experiment (e.g., AR(2) with $K_{\text{MRTS}} = 0, 10, 15, 20$) to test whether temporal flexibility and nonstationary covariates are complementary or redundant. This is currently missing, so the thesis cannot yet quantify interaction effects between the two proposed extensions.

(3) Add replicated uncertainty summaries (mean $\pm$ SD or median with interval) for all MRTS comparisons at each exceedance/quantile level, matching the AR(2) detailed tables, to improve comparability and statistical interpretability.

(4) Add sensitivity analysis for MCMC budget (short vs long chains) and report score stability per unit time, so computational cost comparisons between AR(2), AR(1), NoTS, and MRTS variants are on a fair basis.

4 Ozone data analysis

The real-data analysis applies the model comparison framework to daily maximum 8-hour ozone observations from July 2005. The response is the AQS monitor observation, and the main covariate is the corresponding CMAQ model output. Following the ozone analysis in

Morris et al. (2017), the regression component uses $\mathbf{X}_t(\mathbf{s}) = [1, \text{CMAQ}_t(\mathbf{s})]^\top$, so the model is interpreted as a statistical calibration of CMAQ using site-level observations. The full comparison data set contains 800 monitoring sites observed over 31 days. We use two-fold cross validation, with 400 sites used for training and 400 sites used for validation in each fold.

The purpose of this analysis is not only to reproduce the Morris ozone comparison, but also to evaluate whether the AR(2) time-varying coefficient extension and the MRTS covariate extension are useful for real ozone prediction. A preliminary MERRA-2 analysis with 100 sites and 100 time points showed that AR(2), AR(1), and no-time-series models have very similar MAPE, Brier scores, and quantile scores. The ozone analysis therefore focuses on whether the enlarged model class can improve held-out tail prediction rather than assuming that the more flexible extension must be better.

4.1 Model comparisons

The candidate set contains approximately one hundred model settings. The baseline settings include Gaussian, skew- t , symmetric- t , and max-stable models, with different numbers of knots, censoring thresholds, and time-series choices. For the skew- t and symmetric- t models, the number of knots includes values such as $K = 1, 5, 6, 7, 8, 9, 10$, and 15. Censoring thresholds include no censoring, a moderate threshold such as $T = 50$, and a high threshold such as $T = 75$ ppb. The threshold 75 ppb is scientifically important because it corresponds to the air-quality exceedance level emphasized in the original ozone application.

The extended candidate set adds two families. The first replaces the AR(1) time-varying coefficient structure with the AR(2) structure developed in this thesis. The second adds MRTS covariates to selected Morris-type settings, so that nonstationary spatial structure can enter through multiple-resolution thin-plate spline basis functions. These extensions are compared against the original Morris-type specifications using the same cross-validation splits, so the ranking reflects predictive performance rather than differences in train-test construction.

For each fitted model, posterior predictive samples at the validation sites are used to compute Brier scores and quantile scores. The Brier scores are evaluated at upper-tail empirical thresholds $q(0.90)$, $q(0.95)$, $q(0.98)$, $q(0.99)$, and $q(0.995)$, and quantile scores are evaluated at the same probability levels. Scores are averaged over validation sites, days, and folds. Lower scores indicate better prediction.

Each MCMC chain is run for 30000 iterations with 25000 burn-in iterations. As in the simulation study, convergence of prediction summaries is more important here than direct interpretation of partition-specific latent variables, because multiple-partition models can suffer

from label switching. The model-selection criterion is therefore based on held-out prediction scores, not on a visual comparison of latent partition labels.

4.2 Results

The real-data results support the use of a broader model class than the original baseline alone. In the enlarged ozone comparison, the top two models for Brier and quantile scores at the upper-tail levels include AR(2) and MRTS variants. Tables 3 and 4 summarize the leading models in the primary 400-site training and 400-site validation comparison. At the 0.90 Brier threshold, the best model is an MRTS skew- t specification and the runner-up is the corresponding Morris skew- t specification, with nearly indistinguishable scores. At higher thresholds, the leading model changes across score levels: AR2 models lead the Brier comparisons at 0.95 and 0.995 and the quantile-score comparisons at 0.98, 0.99, and 0.995, whereas Morris skew- t or symmetric- t specifications remain strongest at some intermediate extreme thresholds.

Level	Rank 1 model	Score	Rank 2 model	Score
0.9000	MRTS skew- t , $K = 1$, $T = 0$, AR(1), 10 MRTS covariates	0.04467	Morris skew- t , $K = 1$, $T = 0$, AR(1)	0.04470
0.9500	AR2 skew- t , $K = 7$, $T = 50$	0.02664	AR2 skew- t , $K = 10$, $T = 50$	0.02672
0.9800	Morris skew- t , $K = 6$, $T = 50$, AR(1)	0.01257	Morris skew- t , $K = 5$, $T = 50$, AR(1)	0.01257
0.9900	Morris symmetric- t , $K = 9$, $T = 75$, AR(1)	0.00688	Morris skew- t , $K = 5$, $T = 50$	0.00690
0.9950	AR2 symmetric- t , $K = 15$, $T = 75$	0.00384	Morris symmetric- t , $K = 9$, $T = 75$, AR(1)	0.00385

Table 3: Top two Brier-score models in the US-all ozone validation study. Values in parentheses are held-out scores from the comparison workbook; smaller values indicate better predictive performance.

The ranking is not concentrated on a single universal model. Across Brier scores, the strongest settings typically combine CMAQ, a time-series component, moderate to high thresholds, and a small to moderate number of knots. Across quantile scores, AR(2) settings appear frequently among the best-ranked models, especially for $q(0.95)$ and above. MRTS settings are more selective: they improve some tail exceedance scores, but they are not uniformly preferred across all scores or quantile levels.

The small-sample stability check gives a similar but not identical conclusion. The primary 400/400 split uses empirical thresholds 73.25, 79.75, 88.50, 94.75, and 100.52 for the

Level	Rank 1 model	Score	Rank 2 model	Score
0.9000	MRTS skew- t , $K = 1$, $T = 0$, AR(1), 5 MRTS covariates	5.033	Morris skew- t , $K = 5$, $T = 0$	5.033
0.9500	Morris skew- t , $K = 7$, $T = 50$, AR(1)	2.988	AR2 skew- t , $K = 10$, $T = 0$	2.988
0.9800	AR2 skew- t , $K = 10$, $T = 0$	1.423	Morris skew- t , $K = 9$, $T = 0$	1.424
0.9900	AR2 skew- t , $K = 6$, $T = 50$	0.797	Morris skew- t , $K = 9$, $T = 50$, AR(1)	0.798
0.9950	AR2 skew- t , $K = 6$, $T = 50$	0.441	AR2 skew- t , $K = 8$, $T = 50$	0.442

Table 4: Top two quantile-score models in the US-all ozone validation study. Values in parentheses are held-out scores from the comparison workbook; smaller values indicate better predictive performance.

Level	Rank 1	Rank 2	Rank 3	Rank 4
0.9000	204 (MRTS, 0.041)	55 (Morris's, 0.041)	205 (MRTS, 0.041)	70 (Morris's, 0.041)
0.9500	70 (Morris's, 0.020)	41 (Morris's, 0.020)	12 (Morris's, 0.020)	39 (Morris's, 0.020)
0.9800	204 (MRTS, 0.008)	117 (AR2, 0.008)	67 (Morris's, 0.008)	111 (AR2, 0.008)
0.9900	117 (AR2, 0.003)	15 (Morris's, 0.003)	70 (Morris's, 0.003)	61 (Morris's, 0.003)
0.9950	15 (Morris's, 0.001)	70 (Morris's, 0.001)	61 (Morris's, 0.001)	117 (AR2, 0.001)

Table 5: Validation selection result using 300 sites to predict 100 sites

Level	Rank 1	Rank 2	Rank 3	Rank 4
0.9000	54 (Morris's, 0.050)	105 (AR2, 0.050)	35 (Morris's, 0.050)	215 (MRTS, 0.050)
0.9500	117 (AR2, 0.030)	116 (AR2, 0.030)	114 (AR2, 0.030)	70 (Morris's, 0.030)
0.9800	215 (MRTS, 0.015)	205 (MRTS, 0.015)	111 (AR2, 0.015)	58 (Morris's, 0.015)
0.9900	58 (Morris's, 0.008)	112 (AR2, 0.008)	216 (AR2, 0.008)	67 (Morris's, 0.008)
0.9950	29 (Morris's, 0.004)	35 (Morris's, 0.004)	60 (Morris's, 0.004)	116 (AR2, 0.004)

Table 6: Validation selection result using 200 sites to predict 200 sites

0.90, 0.95, 0.98, 0.99, and 0.995 Brier comparisons, respectively. Under this split, the top setting numbers across these thresholds are 204, 111, 58, 68, and 124. When the experiment is repeated with 300 training sites and 100 validation sites, the leading settings become 204, 70, 204, 117, and 15; with 200 training sites and 200 validation sites, they become 54, 117, 215, 58, and 29. Thus, the same broad model families remain competitive, but the exact best setting changes across thresholds and sample splits. This instability is expected because extreme-ozone validation uses relatively few exceedances, especially at the 0.99 and 0.995 levels. Consequently, the real-data analysis should not be read as identifying a single universally optimal tuning setting, but as evidence that AR(2) and MRTS versions are useful additions to the candidate pool when the target is held-out tail prediction.

Overall, the real-data analysis provides stronger practical motivation for the proposed extensions than the simulation study alone. The AR(2) extension gives additional temporal flexibility and appears in the leading models for several tail quantile scores. The MRTS extension gives a way to enrich the spatial regression component and can improve selected exceedance scores. Therefore, the final modeling recommendation is not to replace the Morris model by a single proposed specification, but to include AR(2) and MRTS versions in the candidate pool and select among them using held-out Brier and quantile scores.

5 MERRA-2 data analysis

This preliminary analysis is conducted on MERRA-2 hourly surface-temperature data to assess whether temporal extensions provide practical predictive gains in a controlled medium-size setting. We use 100 spatial locations and 100 hourly time points, with longitude and latitude retained for each site.

5.1 Data preprocessing and model setup

We randomly split the sites into training and testing sets using seed 123, with 70% of sites used for model fitting and 30% reserved for validation. The regression design is specified as so each time point has an intercept and first-order spatial trend terms. Because the covariates are not standardized in this experiment, we set $\text{nu.upper} = 10$ to stabilize posterior exploration and avoid unstable large- ν behavior during MCMC.

All three temporal specifications are fitted under the same core configuration (skew- t margin, Matérn covariance, $K = 5$ knots, and no censoring threshold) to isolate the effect of tempo-

ral structure. Following the real-data experiment setup, each chain is run for 20000 iterations with 10000 burn-in iterations, yielding 10000 retained posterior draws per model.

5.2 Results

Table 7 summarizes quantile-score (QS) and Brier-score (BS) comparisons for no-temporal, AR(1), and AR(2) specifications. The table shows a clear but subtle trade-off: the no-temporal model performs slightly better on the QS metrics and achieves the smallest BS at the most extreme level (0.99), whereas AR(2) attains the best BS at the 0.90 and 0.95 levels. Because the numerical differences are very small, none of the three temporal structures dominates uniformly; instead, the preferred model depends on whether one emphasizes bulk accuracy, mid-tail exceedances, or the far tail.

The additional summary metrics reinforce this interpretation. MAPE is slightly lower for AR(2) (0.01058699), with AR(1) close behind, while the no-temporal model is somewhat worse on this bulk-accuracy measure. The 95% predictive-interval coverage is also nearly identical across models (approximately 0.912–0.916), which suggests that temporal dependence changes predictive sharpness more than calibration. These modest gains come at a cost: computation time rises from about 777.26 seconds for no-temporal to 1041.65 seconds for AR(1) and 1075.83 seconds for AR(2). Taken together, the results indicate that AR(2) is a reasonable option when slightly better mid-tail behavior is desired, but the improvement is incremental rather than decisive.

Model	QS90	QS95	QS99	BS90	BS95	BS99
No Temporal	4.714	2.522	0.570	0.041	0.018	0.008
AR(1)	4.716	2.522	0.570	0.040	0.016	0.009
AR(2)	4.717	2.522	0.571	0.040	0.016	0.009

Table 7: Comparison of temporal specifications in the MERRA-2 experiment. Smaller values indicate better predictive performance.

6 Discussion

References

- Gneiting, Tilmann and Adrian E Raftery (Mar. 2007). “Strictly Proper Scoring Rules, Prediction, and Estimation”. In: *Journal of the American Statistical Association* 102.477. _eprint: <https://doi.org/10.1198/016214506000001437>, pp. 359–378. ISSN: 0162-1459. DOI: 10.1198/016214506000001437. URL: <https://doi.org/10.1198/016214506000001437> (visited on 11/10/2025).
- Morris, Samuel A. et al. (Sept. 2017). “A Space-Time Skew-t Model for Threshold Exceedances”. In: *Biometrics* 73.3, pp. 749–758. ISSN: 0006-341X. DOI: 10.1111/biom.12644. URL: <https://doi.org/10.1111/biom.12644> (visited on 03/14/2025).
- Tzeng, ShengLi and Hsin-Cheng Huang (Apr. 2018). “Resolution Adaptive Fixed Rank Kriging”. In: *Technometrics* 60.2. _eprint: <https://doi.org/10.1080/00401706.2017.1345701>, pp. 198–208. ISSN: 0040-1706. DOI: 10.1080/00401706.2017.1345701. URL: <https://doi.org/10.1080/00401706.2017.1345701> (visited on 03/22/2026).

