

國立東華大學應用數學系

碩士論文

*Exploration and Extension in Spatial Extremes Incorporating Thin-plate  
Splines Basis Functions*



研究生：謝易軒  
指導教授：黃灝勻 博士

中華民國 115 年 6 月



國立東華大學應用數學系

碩士論文

待定

*To Be Determined*



研究生：謝易軒

指導教授：黃灝勻 博士

中華民國 115 年 6 月



# 學位考試委員會審定書

國立東華大學 \_\_\_\_\_ 系所

研究生 \_\_\_\_\_ 君所提之論文

經本委員會審查並舉行口試，認為

符合碩士學位標準。

學位考試委員會召集人 \_\_\_\_\_ 簽章

指導教授 \_\_\_\_\_ 簽章

委員 \_\_\_\_\_ 簽章

委員 \_\_\_\_\_ 簽章

委員 \_\_\_\_\_ 簽章

系主任 \_\_\_\_\_ 簽章  
(所長)

中華民國 \_\_\_\_\_ 年 \_\_\_\_\_ 月 \_\_\_\_\_ 日



國立東華大學  
NATIONAL DONG HWA UNIVERSITY

學位論文原創性聲明書  
DECLARATION OF THESIS/DISSERTATION ORIGINALITY

學位論文題目：

Thesis/Dissertation Title：

本人在此聲明，所呈交的學位論文是在指導教授黃灝勻的指導下，由個人獨立研究所完成之最終版本。本人對論文內容負責，除了文中已經標註引用處的內容外，論文不包含任何其他他人已經發表或撰寫過的研究成果。對本研究及學位論文做出重要貢獻的個人和組織，均已在文中以明確方式標明。

該論文內容如有違反學術道德或學術規範的行為，如造假、變造、抄襲、研究成果重複發表或未適當引註、以違法或不當手段影響論文審查、不當作者列名等，本人願意承擔由此而產生的法律責任和法律後果。

I declare that the thesis/dissertation herein is the final version of my work, which is composed and accomplished individually under the guidance of my supervisor, Prof. Hao-Yun Huang. I am responsible for the contents of this thesis/dissertation: It contains no research result that was previously published or written by another person. Information derived from published and unpublished work of others has been acknowledged in the text, and a list of references is given. Any contribution made by other individual or organization is explicitly acknowledged in the thesis/dissertation.

If any research misconduct, including fabrication, falsification, or plagiarism in proposing, performing, or reviewing research, or in reporting research results, is discovered in my thesis/dissertation, I am willing to bear corresponding legal responsibilities and all the results therefrom.

聲明人 Declarant：

日期 Date：\_\_\_\_\_ (yyyy/mm/dd)

1140303修訂





## 摘要

為提升美國環境保護署 (US EPA) 對地面臭氧高值的監測能力，Morris et al. (2017) 提出以時空 skew- $t$  過程為基礎的階層式貝氏模型，利用 AQS 監測站資料校正 CMAQ 化學傳輸模式在尾端的預測偏差。該模型結合 Matérn 協方差結構、空間隨機分割、上尾門檻檢查，以及 AR(1) 形式的時變參數，同時刻劃空間非定態性與時間相依性。本論文在此基礎上提出兩項延伸：其一，將時變參數的 AR(1) 先驗推廣為穩態 AR(2) 先驗，並透過 Yule-Walker 關係維持單位邊際變異數，使模型在保有原 skew- $t$  邊際結構的條件下捕捉更長距離的時間相依；其二，將多重解析度薄板樣條 (MRTS) 基底函數納入迴歸成分，以彌補定態 Matérn 協方差難以表現的非定態空間結構。模擬研究與美國臭氧資料分析皆顯示，所提兩項延伸在部分上尾分位下可獲得較原 Morris 模型更低之 Brier 分數，惟其優勢非全面性，而視資料生成機制及目標尾端水準而定。本文因此建議將 AR(2) 與 MRTS 兩種設定納入候選模型集合，並以樣本外預測分數作為最終選模依據。

**關鍵字：**時空 skew- $t$  過程、AR(2) 時變參數、多重解析度薄板樣條、隨機分割、階層式貝氏模型、臭氧預測





## Abstract

To improve the calibration of Community Multiscale Air Quality (CMAQ) model output by ground-level ozone monitoring data, Morris et al. (2017) developed a hierarchical Bayesian framework based on a spatio-temporal skew- $t$  process, combining a Matérn covariance, a random spatial partition, upper-tail censoring, and AR(1) time-varying coefficients to characterize non-stationary spatial dependence and temporal persistence. This thesis extends that framework in two directions. First, the AR(1) prior on the time-varying latent parameters is generalized to a stationary AR(2) prior, with Yule–Walker constraints used to preserve unit marginal variance and hence the spatial skew- $t$  marginal structure at every time point. Second, multi-resolution thin-plate spline (MRTS) basis functions are incorporated into the regression component to capture non-stationary spatial structure that the stationary Matérn covariance cannot represent. Simulation studies and an analysis of U.S. ozone observations show that both extensions can achieve lower Brier scores than the original Morris model at selected upper-tail exceedance levels, although neither extension dominates uniformly across data-generating mechanisms or target quantiles. The proposed extensions therefore enlarge the candidate model class and should be selected among using held-out predictive scores rather than treated as default replacements for the AR(1) baseline.

**Keywords:** spatio-temporal skew- $t$  process; AR(2) time-varying coefficients; multi-resolution thin-plate splines; random partition; hierarchical Bayesian modeling; ozone prediction



# Contents

|                                                        |           |
|--------------------------------------------------------|-----------|
| 學位考試委員會審定書 .....                                       | i         |
| 學位論文原創性聲明書 .....                                       | ii        |
| 摘要 .....                                               | iii       |
| Abstract .....                                         | iv        |
| Contents .....                                         | v         |
| <b>1 Introduction</b> .....                            | <b>1</b>  |
| <b>2 Methodology</b> .....                             | <b>1</b>  |
| 2.1 Spatio-temporal skew-t process .....               | 1         |
| 2.2 Extension to AR(2) time-varying coefficients ..... | 3         |
| 2.3 Incorporate MRTS covariates .....                  | 5         |
| <b>3 Simulation study</b> .....                        | <b>5</b>  |
| 3.1 Design .....                                       | 5         |
| 3.2 Cross validation .....                             | 7         |
| 3.3 Results .....                                      | 8         |
| <b>4 Ozone data analysis</b> .....                     | <b>13</b> |
| 4.1 Model comparisons .....                            | 13        |
| 4.2 Results .....                                      | 14        |
| <b>5 MERRA-2 data analysis</b> .....                   | <b>16</b> |
| 5.1 Data preprocessing and model setup .....           | 16        |
| 5.2 Results .....                                      | 17        |
| <b>6 Discussion</b> .....                              | <b>18</b> |



# 1 Introduction

## 2 Methodology

### 2.1 Spatio-temporal skew-t process

The methodology in this thesis builds on the spatio-temporal skew- $t$  model of Morris et al. (2017). We briefly summarize that model here so that the AR(2) and MRTS extensions in the next two subsections can be stated against a common notation.

Let  $Y_t(\mathbf{s})$  denote the response at spatial location  $\mathbf{s} \in \mathcal{D} \subset \mathcal{R}^2$  on day  $t \in \{1, \dots, n_t\}$ . Following Morris et al. (2017), we model

$$Y_t(\mathbf{s}) = \mathbf{X}_t(\mathbf{s})^\top \boldsymbol{\beta} + \lambda \sigma_t(\mathbf{s}) |z_t(\mathbf{s})| + \sigma_t(\mathbf{s}) v_t(\mathbf{s}), \quad (1)$$

where  $\mathbf{X}_t(\mathbf{s})$  is a  $p$ -vector of covariates,  $\boldsymbol{\beta}$  is the corresponding regression coefficient,  $\lambda \in \mathcal{R}$  is a scalar skewness parameter, and  $v_t(\mathbf{s})$  is a standard Gaussian process with stationary isotropic Matérn correlation

$$\rho(h) = (1 - \gamma) I(h = 0) + \gamma \frac{1}{\Gamma(\nu) 2^{\nu-1}} \left( \sqrt{2\nu} \frac{h}{\varphi} \right)^\nu K_\nu \left( \sqrt{2\nu} \frac{h}{\varphi} \right), \quad (2)$$

where  $\varphi > 0$  is the spatial range,  $\nu > 0$  is the smoothness,  $\gamma \in [0, 1]$  is the proportion of variance attributed to the spatial component,  $K_\nu$  is the modified Bessel function of the second kind, and  $h = \|\mathbf{s}_1 - \mathbf{s}_2\|$ . Marginalizing over  $|z_t(\mathbf{s})|$  and  $\sigma_t(\mathbf{s})$ , the finite-dimensional vector  $\mathbf{Y}_t = [Y_t(\mathbf{s}_1), \dots, Y_t(\mathbf{s}_n)]^\top$  follows an  $n$ -dimensional skew- $t$  distribution; consequently (1) provides asymmetric and heavy-tailed margins together with asymptotic dependence between nearby sites.

To remove the long-range asymptotic dependence that arises when a single scalar  $z$  and  $\sigma^2$  are shared across the entire domain, the model uses a daily random partition. For each day  $t$ , let  $\mathbf{w}_{t1}, \dots, \mathbf{w}_{tK}$  denote  $K$  spatial knots in  $\mathcal{D}$ , and define the partition subregions

$$P_{tk} = \{\mathbf{s} \in \mathcal{D} : k = \arg \min_\ell \|\mathbf{s} - \mathbf{w}_{t\ell}\|\}, \quad k = 1, \dots, K. \quad (3)$$

For  $\mathbf{s} \in P_{tk}$  we set  $z_t(\mathbf{s}) = z_{tk}$  and  $\sigma_t^2(\mathbf{s}) = \sigma_{tk}^2$ , with  $z_{tk} \stackrel{iid}{\sim} N(0, 1)$  and  $\sigma_{tk}^2 \stackrel{iid}{\sim} \text{IG}(a/2, b/2)$ . Within each subregion the model retains the skew- $t$  form, while sites in different partitions are not asymptotically dependent.

Because the EPA compliance criterion focuses on exceedances of a high level  $L$ , the model fits only the upper tail by censoring the data at a pre-specified threshold  $T \leq L$ . Writing the partially censored response  $\tilde{Y}_t(\mathbf{s}) = \max\{Y_t(\mathbf{s}), T\}$ , inference is based on  $\tilde{\mathbf{Y}}_t = [\tilde{Y}_t(\mathbf{s}_1), \dots, \tilde{Y}_t(\mathbf{s}_n)]^\top$ , and the censored values below  $T$  are imputed inside the MCMC.

Daily temporal dependence is introduced on the latent parameters that drive the tail, namely the knot locations  $\mathbf{w}_{tk}$ , the skewness latents  $z_{tk}$ , and the scale parameters  $\sigma_{tk}^2$ . To preserve the marginal skew- $t$  distribution at every day, each parameter is first transformed to a standard normal scale via the probability integral transform:

$$w_{tki}^* = \Phi^{-1} \left[ \frac{w_{tki} - \min(\mathbf{s}_i)}{\max(\mathbf{s}_i) - \min(\mathbf{s}_i)} \right], \quad i = 1, 2, \quad (4)$$

$$z_{tk}^* = \Phi^{-1} \{ \text{HN}(z_{tk}) \}, \quad (5)$$

$$\sigma_{tk}^{2*} = \Phi^{-1} \{ \text{IG}(\sigma_{tk}^2) \}, \quad (6)$$

where  $\Phi$  is the standard normal distribution function, HN is the half-normal distribution function, and IG is the inverse-gamma distribution function with parameters  $(a/2, b/2)$ . After transformation, each  $\mathbf{w}_{tk}^*, z_{tk}^*, \sigma_{tk}^{2*}$  has a standard normal marginal distribution, and Morris et al. (2017) place a stationary AR(1) Gaussian time series on each of them:

$$\mathbf{w}_{tk}^* \mid \mathbf{w}_{t-1,k}^* \sim N_2(\phi_w \mathbf{w}_{t-1,k}^*, (1 - \phi_w^2) \mathbf{I}_2), \quad (7)$$

$$z_{tk}^* \mid z_{t-1,k}^* \sim N(\phi_z z_{t-1,k}^*, 1 - \phi_z^2), \quad (8)$$

$$\sigma_{tk}^{2*} \mid \sigma_{t-1,k}^{2*} \sim N(\phi_\sigma \sigma_{t-1,k}^{2*}, 1 - \phi_\sigma^2), \quad (9)$$

with  $|\phi_w|, |\phi_z|, |\phi_\sigma| < 1$ , and initial values  $\mathbf{w}_{1k}^* \sim N_2(\mathbf{0}, \mathbf{I}_2)$ ,  $z_{1k}^* \sim N(0, 1)$ ,  $\sigma_{1k}^{2*} \sim N(0, 1)$ . Transforming back through (4)–(6) gives  $\mathbf{w}_{tk} \sim \text{Unif}(\mathcal{D})$ ,  $z_{tk} \sim \text{HN}(0, 1)$ , and  $\sigma_{tk}^2 \sim \text{IG}(a/2, b/2)$  at every  $t$ , so that the spatial skew- $t$  structure of (1) is preserved at each day. The complete hierarchical model is fitted via MCMC with Gibbs and Metropolis–Hastings steps, with predictions at unobserved sites obtained by Bayesian kriging on the multivariate normal full conditional of  $\mathbf{Y}_t$ .

This skew- $t$  baseline – partitioned, censored, and AR(1)-correlated in time – defines the reference specification against which the two proposed extensions of this thesis are evaluated.



## 2.2 Extension to AR(2) time-varying coefficients

The AR(1) time series in (7)–(9) links each transformed parameter on day  $t$  to its value on day  $t - 1$  only. This may be too restrictive for daily atmospheric variables such as ozone, where persistence and weather-driven cycles can carry information across multiple lags. We therefore replace the AR(1) prior on each transformed parameter by a stationary AR(2) prior, while preserving the unit marginal variance required by the probability integral transforms in (4)–(6).

**The AR(2) recursion.** For a generic transformed parameter  $\{Y_t\}_{t=1}^{n_t}$  – standing in for any of the scalar components of  $\mathbf{w}_{tk}^*$ , or for  $z_{tk}^*$  or  $\sigma_{tk}^{2*}$  at a fixed knot  $k$  – the AR(2) prior is

$$Y_t = \phi_1 Y_{t-1} + \phi_2 Y_{t-2} + \epsilon_t, \quad \epsilon_t \stackrel{iid}{\sim} N(0, \sigma_\epsilon^2), \quad t \geq 3. \quad (10)$$

For the recursion to define a stationary process,  $(\phi_1, \phi_2)$  must lie inside the stationarity triangle

$$|\phi_2| < 1, \quad \phi_1 + \phi_2 < 1, \quad \phi_2 - \phi_1 < 1. \quad (11)$$

Just as in the AR(1) case, the AR(2) parameters  $(\phi_1^{(w)}, \phi_2^{(w)})$ ,  $(\phi_1^{(z)}, \phi_2^{(z)})$ , and  $(\phi_1^{(\sigma)}, \phi_2^{(\sigma)})$  are separately defined for  $\mathbf{w}_{tk}^*$ ,  $z_{tk}^*$ , and  $\sigma_{tk}^{2*}$ .

**Yule–Walker moments under unit marginal variance.** We require  $\gamma_0 := \text{Var}(Y_t) = 1$  so that the inverse PIT applied to  $Y_t$  recovers the half-normal, inverse-gamma, or uniform target margins on the original scale. Multiplying (10) by  $Y_{t-h}$  and taking expectations yields the Yule–Walker recursion

$$\gamma_h = \phi_1 \gamma_{h-1} + \phi_2 \gamma_{h-2}. \quad (12)$$

Setting  $h = 1$  and using the symmetry  $\gamma_{-1} = \gamma_1$  gives

$$\gamma_1 = \frac{\phi_1}{1 - \phi_2}, \quad \gamma_2 = \phi_1 \gamma_1 + \phi_2, \quad (13)$$

and multiplying (10) by  $Y_t$  pins down the innovation variance

$$\sigma_\epsilon^2 = 1 - \phi_1 \gamma_1 - \phi_2 \gamma_2. \quad (14)$$

Setting  $\phi_2 = 0$  in (13)–(14) recovers the AR(1) form  $\gamma_1 = \phi_1$ ,  $\sigma_\epsilon^2 = 1 - \phi_1^2$ , so the AR(2) prior strictly contains the AR(1) prior used by Morris et al. (2017).

**Stationary joint of the initial pair.** Because the recursion (10) requires two seeds, the joint distribution of  $(Y_1, Y_2)$  has to be specified separately. Independent  $N(0, 1)$  initial values would create a transient regime in which  $\text{Var}(Y_t) \neq 1$  for the first several days, breaking the PIT mechanism that produces the correct skew- $t$  margin. Stationarity therefore forces the bivariate Gaussian initial condition

$$\begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix} \sim N_2 \left( \begin{pmatrix} 0 \\ 0 \end{pmatrix}, \begin{pmatrix} 1 & \gamma_1 \\ \gamma_1 & 1 \end{pmatrix} \right), \quad (15)$$

which yields the conditional density used when updating  $Y_2$  alone:

$$Y_2 | Y_1 \sim N(\gamma_1 Y_1, 1 - \gamma_1^2). \quad (16)$$

Note that the mean of (16) is  $\gamma_1 Y_1$ , not  $\phi_1 Y_1$ , and its variance is  $1 - \gamma_1^2$ , not  $\sigma_\epsilon^2$ ; the two expressions coincide only when  $\phi_2 = 0$ . In the AR(2) implementation it is essential that both the simulator and the MCMC density at  $t = 2$  use (16) so that the marginal property  $\text{Var}(Y_t) = 1$  holds at every  $t$ .

**Full conditional with three prior factors.** The AR(2) Markov structure factorizes the joint prior of  $(Y_1, \dots, Y_{n_t})$  as

$$p(Y_1, \dots, Y_{n_t}) = p(Y_1, Y_2) \prod_{s=3}^{n_t} p(Y_s | Y_{s-1}, Y_{s-2}), \quad (17)$$

with  $p(Y_1, Y_2)$  given by (15) and each conditional factor for  $s \geq 3$  following (10). For an interior time point  $3 \leq t \leq n_t - 2$ , the value  $Y_t$  enters (17) through three factors: the self factor  $p(Y_t | Y_{t-1}, Y_{t-2})$ , the lag-1 factor  $p(Y_{t+1} | Y_t, Y_{t-1})$ , and the lag-2 factor  $p(Y_{t+2} | Y_{t+1}, Y_t)$ . The full conditional in the Gibbs-within-MH update of  $Y_t$  is therefore

$$\pi(Y_t | Y_{-t}, \mathbf{Y}) \propto L(\mathbf{Y} | Y_t) \cdot \underbrace{p(Y_t | Y_{t-1}, Y_{t-2})}_{\text{self}} \cdot \underbrace{p(Y_{t+1} | Y_t, Y_{t-1})}_{\text{lag-1}} \cdot \underbrace{p(Y_{t+2} | Y_{t+1}, Y_t)}_{\text{lag-2}}, \quad (18)$$

where  $L(\mathbf{Y} \mid Y_t)$  is the spatial likelihood contributed by  $Y_t$  through the partition and Gaussian process structure of (1). The lag-2 factor is absent for  $t = n_t - 1$  and both lag-1 and lag-2 factors are absent for  $t = n_t$ . Under a symmetric Gaussian random-walk proposal  $Y_t^{(c)} \sim N(Y_t, s^2)$ , the Metropolis–Hastings acceptance ratio is

$$R = \min \left( 1, \frac{\pi(Y_t^{(c)} \mid Y_{-t}, \mathbf{Y})}{\pi(Y_t \mid Y_{-t}, \mathbf{Y})} \right), \quad (19)$$

in which all three prior factors must appear; omitting the lag-2 factor distorts the stationary distribution of the chain away from the AR(2) joint and produces a lag-2 autocorrelation that does not match the Yule–Walker value in (13).

**Application to all three transformed parameters.** The AR(2) prior (10) together with the stationary initial condition (15) replaces (7)–(9) for each transformed parameter:

$$\mathbf{w}_{tk}^* \mid \mathbf{w}_{t-1,k}^*, \mathbf{w}_{t-2,k}^* \sim N_2 \left( \phi_1^{(w)} \mathbf{w}_{t-1,k}^* + \phi_2^{(w)} \mathbf{w}_{t-2,k}^*, \sigma_{\epsilon,w}^2 \mathbf{I}_2 \right), \quad (20)$$

$$z_{tk}^* \mid z_{t-1,k}^*, z_{t-2,k}^* \sim N \left( \phi_1^{(z)} z_{t-1,k}^* + \phi_2^{(z)} z_{t-2,k}^*, \sigma_{\epsilon,z}^2 \right), \quad (21)$$

$$\sigma_{tk}^{2*} \mid \sigma_{t-1,k}^{2*}, \sigma_{t-2,k}^{2*} \sim N \left( \phi_1^{(\sigma)} \sigma_{t-1,k}^{2*} + \phi_2^{(\sigma)} \sigma_{t-2,k}^{2*}, \sigma_{\epsilon,\sigma}^2 \right), \quad (22)$$

where each pair  $(\phi_1^{(\cdot)}, \phi_2^{(\cdot)})$  lies inside the stationarity triangle (11), the innovation variances are obtained from (13)–(14), and the initial pairs  $(\mathbf{w}_{1k}^*, \mathbf{w}_{2k}^*)$ ,  $(z_{1k}^*, z_{2k}^*)$ ,  $(\sigma_{1k}^{2*}, \sigma_{2k}^{2*})$  follow the bivariate Gaussian (15) component-wise. After inverse transformation through (4)–(6), the back-transformed parameters  $\mathbf{w}_{tk}$ ,  $z_{tk}$ ,  $\sigma_{tk}^2$  retain the same marginal distributions as in the AR(1) baseline, so that the spatial skew- $t$  model (1) is preserved at every  $t$ . Setting  $\phi_2^{(\cdot)} = 0$  in (20)–(22) recovers the Morris AR(1) specification, so the AR(2) extension nests the original time-series model and adds one additional autoregressive coefficient per latent process.

## 2.3 Incorporate MRTS covariates

# 3 Simulation study

## 3.1 Design

The simulation study is used to evaluate whether the proposed extensions improve predictive performance beyond the space-time skew- $t$  threshold model of Morris et al. (2017). The

benchmark simulation follows the structure of Morris et al.: observations are generated at  $n_s = 144$  spatial sites, sampled uniformly on  $[0, 10] \times [0, 10]$ , over  $n_t = 50$  independent time points. The mean structure is constant in the data generation,  $\mathbf{X}^\top \boldsymbol{\beta} = 10$ , while the fitted models include an intercept and first-order spatial trend. The main data-generating settings include Gaussian data, skew- $t$  data with one knot, skew- $t$  data with five knots, a max-stable setting, and a Brown-Resnick max-stable setting. For each setting, 50 simulated data sets are generated.

For the baseline comparison, five Bayesian models are considered:

1. Gaussian marginal distribution with  $K = 1$  knot;
2. Skew- $t$  marginal distribution with  $K = 1$  knot and no thresholding;
3. Symmetric- $t$  marginal distribution with  $K = 1$  knot and threshold  $T = q(0.80)$ ;
4. Skew- $t$  marginal distribution with  $K = 5$  knots and no thresholding;
5. Symmetric- $t$  marginal distribution with  $K = 5$  knots and threshold  $T = q(0.80)$ .

Here  $q(0.80)$  denotes the empirical 80th percentile of the simulated response. These five models represent the main combinations of tail behavior, thresholding, and spatial partitioning used in the Morris model.

Two additional simulation experiments are conducted for the extensions in this thesis. First, to evaluate the AR(2) time-varying coefficient structure, data are generated with temporal dependence in the latent  $\tau$ ,  $\mathbf{w}$ , and  $\mathbf{z}$  processes. The experiment uses  $n = 100$  sites and  $T = 30$  time points, with 90 sites used for training and 10 sites withheld for prediction. We repeat the experiment 20 times and compare three temporal structures: no time series, AR(1), and AR(2). Each fitted model uses a skew- $t$  marginal distribution with  $K = 5$  knots and no thresholding, and the MCMC is run for 5000 iterations with 2500 burn-in iterations.

A complementary AR(2) experiment embeds the temporal mechanism directly into the benchmark simulation, so that the AR(2) analysis model can be tested against data with a known autoregressive structure. Two further data-generating settings are added. Both produce one-knot skew- $t$  data ( $K = 1$ ,  $\lambda = 3$ ) on the same  $n_s = 144$  sites and  $n_t = 50$  time points as the benchmark, but the latent  $\tau$ ,  $\mathbf{w}$ , and  $\mathbf{z}$  processes now evolve through a fixed AR(2) recursion (20)–(22) whose coefficients are held constant across time. Setting 9 uses a strongly dependent coefficient pair  $(\phi_1, \phi_2) = (0.80, -0.35)$ , and Setting 10 a weakly dependent pair  $(\phi_1, \phi_2) = (0.12, -0.05)$ ; both pairs lie inside the stationarity triangle (11). For each

setting, 50 data sets are simulated and split into 100 training and 44 validation sites. Each data set is fit by seven analysis methods: the five baseline models listed above, together with two AR(2)-augmented skew- $t$  models—one with  $K = 1$  knot and one with  $K = 5$  knots—whose time-varying coefficients follow the AR(2) prior (20)–(22). Because the data are themselves generated with AR(2) dependence, this comparison gives the temporal extension its most favorable test.

Second, to evaluate the MRTS covariate extension, we augment the original five models with multi-resolution thin-plate spline basis functions (Tzeng and Huang, 2018). These basis functions are added to the regression component in order to capture nonstationary spatial structure that may not be fully represented by a stationary Matérn covariance. We focus on simulation settings 4 and 5: setting 4 is a one-knot skew- $t$  data-generating process with  $\lambda = 3$ , and setting 5 is a five-knot partitioned skew- $t$  data-generating process with  $\lambda = 3$ . For setting 4, each original model is compared with versions using 10, 15, and 20 MRTS covariates. For setting 5, each original model is compared with a version using 15 MRTS covariates. Each setting uses 50 simulated data sets, and each MCMC chain is run for 20000 iterations with 10000 burn-in iterations.

### 3.2 Cross validation

Predictive performance is evaluated by withholding spatial sites from model fitting and scoring posterior predictive distributions at the held-out sites. In the benchmark simulation, 100 sites are used as training sites and 44 sites are withheld for validation. In the AR(2) experiment, the split is 90 training sites and 10 testing sites. The MRTS experiment follows the benchmark simulation structure and compares each MRTS-augmented model with the corresponding non-MRTS model on the same simulated data sets.


Because the main inferential target is prediction of high-level exceedances, the primary scoring rule is the Brier score (Gneiting and Raftery, 2007). For an exceedance level  $L$ , the Brier score is  $BS(L) = \{I(y > L) - \hat{P}(Y > L)\}^2$ , where  $y$  is the held-out observation and  $\hat{P}(Y > L)$  is the posterior predictive probability of exceeding  $L$ . The score is averaged over all held-out sites and time points, and lower values indicate better exceedance prediction.

For the extension experiments, we also record MAPE and quantile scores. The quantile score at level  $\alpha$  is used to evaluate the accuracy of posterior predictive quantiles, with lower values again indicating better prediction. The reported exceedance and quantile levels are concentrated in the upper tail: the AR(2) experiment uses  $q(0.90)$ ,  $q(0.95)$ ,  $q(0.98)$ , and  $q(0.99)$ , while the MRTS experiment uses levels from  $q(0.90)$  to  $q(0.995)$ .

### 3.3 Results

The benchmark simulation supports the main conclusion of Morris et al. (2017): when the fitted model matches the data-generating mechanism, it tends to improve predictive performance, and when the data are non-Gaussian the skew- $t$  models generally improve over a Gaussian model. Thresholded models are not uniformly better in the simulation setting, which is expected because most simulated data are generated directly from the full skew- $t$  model rather than from a censored tail model. The simulation therefore confirms that the skew- $t$  framework provides flexibility for asymmetric and heavy-tailed data while remaining competitive with Gaussian methods when the data are light-tailed and symmetric.

Figure 1 examines the AR(2) extension directly. All seven analysis methods are fitted to the two fixed-AR(2) data-generating settings, and for each method the Brier score is plotted relative to the Gaussian model across upper-tail threshold quantiles. A value below the dashed reference line at 1 means the method predicts exceedances better than the Gaussian baseline at that quantile, and a value above it means worse.



(a) Setting 9: strong AR(2) dependence,  $(\phi_1, \phi_2) = (0.80, -0.35)$ . (b) Setting 10: weak AR(2) dependence,  $(\phi_1, \phi_2) = (0.12, -0.05)$ .

Figure 1: Relative Brier score (versus the Gaussian model) by threshold quantile for the seven analysis methods fitted to the fixed-AR(2) data-generating settings. Values below the dashed reference line at 1 indicate better exceedance prediction than the Gaussian baseline; values above it indicate worse.

Under strong serial dependence (Setting 9, Figure 1a), every skew- $t$  and symmetric- $t$  method except the five-knot thresholded  $t$ -model improves on the Gaussian baseline, with relative Brier scores between roughly 0.94 and 0.98. The one-knot skew- $t$  model, which matches the  $K = 1$  data-generating truth, is the best method, with a relative score near 0.95. Critically, adding the AR(2) temporal structure does not help: the one-knot AR(2) skew- $t$  curve

sits essentially on top of—and slightly above—the one-knot non-temporal skew- $t$  curve, and the five-knot AR(2) model tracks its non-temporal counterpart. The five-knot thresholded  $t$ -model is the only method worse than Gaussian, and its relative score degrades sharply at the 0.99 quantile.

Under weak serial dependence (Setting 10, Figure 1b) the pattern shifts. The five-knot skew- $t$  and five-knot thresholded  $t$  models are now clearly best, with relative Brier scores mostly between about 0.88 and 0.94, while the one-knot models—including the one-knot AR(2) model—are essentially indistinguishable from the Gaussian baseline. Once again the AR(2) and non-temporal curves nearly coincide at each knot count. Far-tail estimates are noisy: at the 0.995 quantile several methods rise above 1 because few observations exceed such extreme thresholds.

Two messages emerge. First, on data that are explicitly generated with AR(2) dependence—the most favorable possible setting for the temporal extension—fitting an AR(2) analysis model yields no measurable Brier-score gain over the corresponding non-temporal skew- $t$  model at any quantile. Predictive accuracy is governed by the marginal and spatial flexibility of the model, namely the skew- $t$  margin and the number of knots, rather than by the autoregressive order of the latent coefficients. Second, the preferred knot count depends on the dependence regime: strong temporal dependence favors the parsimonious one-knot model that matches the generating truth, whereas weak dependence favors the richer five-knot partition. This reinforces the conclusion that the AR(2) structure should be treated as a candidate option rather than a default improvement.

A complementary set of missing-at-location experiments places the AR(2), AR(1), and no-time-series specifications against a range of dependence regimes; the resulting Brier scores at the 0.90, 0.95, and 0.99 exceedance levels are collected in Table 1. AR(2) improves mid-tail Brier scores in some scenarios – most visibly when one or more of the latent autoregressive coefficients is large – but AR(1) or the no-time-series baseline remains preferable in other scenarios or at specific tail levels, and the gains over AR(1) are typically smaller than the gap between any time-series model and the no-time-series baseline in runtime. Temporal flexibility therefore does not translate into uniform predictive dominance.

The MRTS evidence is qualitatively similar: in settings 4 and 5, augmenting the regression component with MRTS basis functions can improve selected upper-tail Brier or quantile scores, but the gains are not stable across model families or quantile levels, and computation time grows visibly with  $K_{\text{MRTS}}$ , most clearly at  $K_{\text{MRTS}} = 20$ .

Figures 2 and 3 make this explicit by comparing methods 2–5 (skew- $t$  and symmetric- $t$

| Scenario (missing-at-location) | BS90 (AR2/AR1/NoTS)  | BS95 (AR2/AR1/NoTS)  | BS99 (AR2/AR1/NoTS)  | Time(s)<br>(AR2/AR1/NoTS) |
|--------------------------------|----------------------|----------------------|----------------------|---------------------------|
| $\sigma^2$ AR(2) only          | 0.0129/0.0130/0.0128 | 0.0106/0.0106/0.0106 | 0.0032/0.0032/0.0032 | 54.78/54.53/51.41         |
| All large $\phi$               | 0.0344/0.0346/0.0352 | 0.0269/0.0267/0.0268 | 0.0071/0.0061/0.0069 | 138.47/135.74/109.64      |
| All small $\phi$               | 0.0267/0.0267/0.0266 | 0.0178/0.0178/0.0179 | 0.0059/0.0060/0.0060 | 232.74/228.68/187.20      |
| Big $\phi_\tau$ , others small | 0.0388/0.0390/0.0389 | 0.0181/0.0182/0.0182 | 0.0045/0.0044/0.0043 | 135.17/134.59/102.93      |

Table 1: Selected AR(2) simulation results under missing-at-location settings; lower values indicate better predictive performance.

with  $K = 1$  or  $K = 5$  knots) at high threshold quantiles ( $0.90 \leq q \leq 0.995$ ), with each figure showing two side-by-side panels: the left panel is the no-MRTS baseline and the right panel adds an MRTS basis of rank  $K_{\text{MRTS}} = 15$ . All scores are normalized by the Gaussian baseline (method 1) computed without MRTS, so values below the dashed reference line at 1 indicate better exceedance prediction than the Gaussian baseline. Setting 4 (skew- $t$  with  $K = 1$  knot,  $\lambda = 3$ ) is shown in Figure 2, and setting 5 (skew- $t$  with  $K = 5$  knots,  $\lambda = 3$ ) in Figure 3.




Figure 2: Setting 4 (data-generating process: skew- $t$  with  $K = 1$  knot,  $\lambda = 3$ ). Relative Brier scores of methods 2–5 versus the Gaussian baseline at threshold quantiles  $0.90 \leq q \leq 0.995$ , with  $K_{\text{MRTS}} = 15$ . The left panel shows the no-MRTS baseline and the right panel shows the same methods with the MRTS basis added. The matching method (skew- $t$  with  $K = 1$ , method 2) is best at every quantile, the over-specified five-knot threshold- $t$  method (method 5) is worse than Gaussian, and the right panel is nearly indistinguishable from the left.

Two consistent patterns emerge. First, the data-generating knot count is the dominant axis: in setting 4 the one-knot skew- $t$  method (method 2) is best at every quantile, and in






Figure 3: Setting 5 (data-generating process: skew- $t$  with  $K = 5$  knots,  $\lambda = 3$ ). Relative Brier scores of methods 2–5 versus the Gaussian baseline; left panel is the no-MRTS baseline, right panel adds an MRTS basis of rank  $K_{\text{MRTS}} = 15$ . The matching five-knot skew- $t$  method (method 4) is decisively the best across all quantiles, with relative Brier scores between roughly 0.84 and 0.92. All four methods improve on the Gaussian baseline, in contrast with setting 4 where the over-specified method 5 was worse than Gaussian. The right panel mirrors the left panel closely, indicating that the MRTS basis adds little beyond what the spatial skew- $t$  structure already captures.

setting 5 the five-knot skew- $t$  method (method 4) is best at every quantile. Mis-specifying  $K$  on the under-side (fitting  $K = 1$  to a  $K = 5$  truth, as method 2 does in setting 5) gives only marginal gains over Gaussian, whereas mis-specifying  $K$  on the over-side (fitting  $K = 5$  to a  $K = 1$  truth, as method 4 does in setting 4) is competitive but not optimal. Second, the MRTS basis is essentially flat across the configurations studied here: in both settings the right (with-MRTS) panel reproduces the left (without-MRTS) panel almost exactly, and the same pattern was observed at the other ranks examined ( $K_{\text{MRTS}} \in \{10, 20\}$ , results not shown). The MRTS basis therefore does not change the within-setting ranking, and its added computational cost is not repaid by tail-score improvement under these data-generating processes. Table 2 places AR(2) and MRTS side-by-side as proposed extensions in a common decision summary.

| Extension                       | Simulation setup                                                                                                                        | Main empirical pattern                                                                                                                                                        | Practical implication                                                                          |
|---------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| AR(2) time-varying coefficients | Missing-at-location experiments with AR(2), AR(1), and no-time-series fitted to multiple dependence regimes                             | AR(2) can improve selected mid-tail Brier scores, but runtime advantages are often with AR(1) or no-time-series; no uniform winner across regimes                             | Use AR(2) as a candidate temporal model, not as a default replacement                          |
| MRTS covariates                 | Benchmark simulation settings 4–5 with additional multi-resolution thin-plate spline covariates ( $K_{\text{MRTS}} = 10, 15, 20$ or 15) | Improvements appear at selected tail levels, but average differences are small and not stable across quantiles/model classes; runtime increases with larger $K_{\text{MRTS}}$ | Use MRTS selectively when nonstationarity is plausible and tail-score gains justify added cost |

Table 2: Integrated simulation summary for the two proposed extensions.

Overall, the simulation study indicates that both extensions enlarge the usable model class but that neither uniformly dominates the Morris baseline. The AR(2) prior provides temporal flexibility and the MRTS covariates enrich the spatial regression component, yet their empirical benefits depend on the data-generating mechanism and on the target exceedance level. In the ozone analysis that follows, AR(2) and MRTS specifications are therefore treated as candidate options to be selected by held-out predictive scores rather than as default replacements for the AR(1) Morris model.

## 4 Ozone data analysis

The real-data analysis applies the model comparison framework to daily maximum 8-hour ozone observations from July 2005. The response is the AQS monitor observation, and the main covariate is the corresponding CMAQ model output. Following the ozone analysis in Morris et al. (2017), the regression component uses  $\mathbf{X}_t(\mathbf{s}) = [1, \text{CMAQ}_t(\mathbf{s})]^\top$ , so the model is interpreted as a statistical calibration of CMAQ using site-level observations. The full

comparison data set contains 800 monitoring sites observed over 31 days. We use two-fold cross validation, with 400 sites used for training and 400 sites used for validation in each fold.

The purpose of this analysis is not only to reproduce the Morris ozone comparison, but also to evaluate whether the AR(2) time-varying coefficient extension and the MRTS covariate extension are useful for real ozone prediction. A preliminary MERRA-2 analysis with 100 sites and 100 time points showed that AR(2), AR(1), and no-time-series models have very similar MAPE, Brier scores, and quantile scores. The ozone analysis therefore focuses on whether the enlarged model class can improve held-out tail prediction rather than assuming that the more flexible extension must be better.

## 4.1 Model comparisons

The candidate set contains approximately one hundred model settings. The baseline settings include Gaussian, skew- $t$ , symmetric- $t$ , and max-stable models, with different numbers of knots, censoring thresholds, and time-series choices. For the skew- $t$  and symmetric- $t$  models, the number of knots includes values such as  $K = 1, 5, 6, 7, 8, 9, 10$ , and 15. Censoring thresholds include no censoring, a moderate threshold such as  $T = 50$ , and a high threshold such as  $T = 75$  ppb. The threshold 75 ppb is scientifically important because it corresponds to the air-quality exceedance level emphasized in the original ozone application.

The extended candidate set adds two families. The first replaces the AR(1) time-varying coefficient structure with the AR(2) structure developed in this thesis. The second adds MRTS covariates to selected Morris-type settings, so that nonstationary spatial structure can enter through multiple-resolution thin-plate spline basis functions. These extensions are compared against the original Morris-type specifications using the same cross-validation splits, so the ranking reflects predictive performance rather than differences in train-test construction.

For each fitted model, posterior predictive samples at the validation sites are used to compute Brier scores and quantile scores. The Brier scores are evaluated at upper-tail empirical thresholds  $q(0.90)$ ,  $q(0.95)$ ,  $q(0.98)$ ,  $q(0.99)$ , and  $q(0.995)$ , and quantile scores are evaluated at the same probability levels. Scores are averaged over validation sites, days, and folds. Lower scores indicate better prediction.

Each MCMC chain is run for 30000 iterations with 25000 burn-in iterations. As in the simulation study, convergence of prediction summaries is more important here than direct interpretation of partition-specific latent variables, because multiple-partition models can suffer from label switching. The model-selection criterion is therefore based on held-out prediction scores, not on a visual comparison of latent partition labels.

## 4.2 Results

The real-data results support enlarging the candidate model class beyond the Morris baseline. Tables 3 and 4 list the top two models, by Brier score and by quantile score respectively, in the primary 400-train/400-validation comparison. AR(2) and MRTS variants appear repeatedly among the leading specifications: at the 0.90 Brier level the best model is an MRTS skew- $t$  specification, only marginally ahead of the corresponding Morris specification; AR(2) models lead the Brier comparisons at 0.95 and 0.995 and the quantile-score comparisons at 0.98, 0.99, and 0.995. At intermediate exceedance levels, however, Morris skew- $t$  or symmetric- $t$  specifications remain strongest, so no single proposed extension dominates the entire tail.

| Level  | Rank 1 model                                                             | Score   | Rank 2 model                                          | Score   |
|--------|--------------------------------------------------------------------------|---------|-------------------------------------------------------|---------|
| 0.9000 | <b>MRTS</b> skew- $t$ , $K = 1$ , $T = 0$ ,<br>AR(1), 10 MRTS covariates | 0.04467 | Morris skew- $t$ , $K = 1$ , $T = 0$ , AR(1)          | 0.04470 |
| 0.9500 | <b>AR2</b> skew- $t$ , $K = 7$ , $T = 50$                                | 0.02664 | <b>AR2</b> skew- $t$ , $K = 10$ , $T = 50$            | 0.02672 |
| 0.9800 | Morris skew- $t$ , $K = 6$ , $T = 50$ ,<br>AR(1)                         | 0.01257 | Morris skew- $t$ , $K = 5$ , $T = 50$ ,<br>AR(1)      | 0.01257 |
| 0.9900 | Morris symmetric- $t$ , $K = 9$ , $T = 75$ ,<br>AR(1)                    | 0.00688 | Morris skew- $t$ , $K = 5$ , $T = 50$                 | 0.00690 |
| 0.9950 | <b>AR2</b> symmetric- $t$ , $K = 15$ , $T = 75$                          | 0.00384 | Morris symmetric- $t$ , $K = 9$ , $T = 75$ ,<br>AR(1) | 0.00385 |

Table 3: Top two Brier-score models in the US-all ozone validation study. Values in parentheses are held-out scores from the comparison workbook; smaller values indicate better predictive performance.

| Level  | Rank 1 model                                                            | Score | Rank 2 model                                     | Score |
|--------|-------------------------------------------------------------------------|-------|--------------------------------------------------|-------|
| 0.9000 | <b>MRTS</b> skew- $t$ , $K = 1$ , $T = 0$ ,<br>AR(1), 5 MRTS covariates | 5.033 | Morris skew- $t$ , $K = 5$ , $T = 0$             | 5.033 |
| 0.9500 | Morris skew- $t$ , $K = 7$ , $T = 50$ ,<br>AR(1)                        | 2.988 | <b>AR2</b> skew- $t$ , $K = 10$ , $T = 0$        | 2.988 |
| 0.9800 | <b>AR2</b> skew- $t$ , $K = 10$ , $T = 0$                               | 1.423 | Morris skew- $t$ , $K = 9$ , $T = 0$             | 1.424 |
| 0.9900 | <b>AR2</b> skew- $t$ , $K = 6$ , $T = 50$                               | 0.797 | Morris skew- $t$ , $K = 9$ , $T = 50$ ,<br>AR(1) | 0.798 |
| 0.9950 | <b>AR2</b> skew- $t$ , $K = 6$ , $T = 50$                               | 0.441 | <b>AR2</b> skew- $t$ , $K = 8$ , $T = 50$        | 0.442 |

Table 4: Top two quantile-score models in the US-all ozone validation study. Values in parentheses are held-out scores from the comparison workbook; smaller values indicate better predictive performance.

The ranking does not concentrate on a single universally optimal specification. Among the

Brier-score leaders, the strongest settings typically combine CMAQ, a time-series component, a moderate-to-high censoring threshold, and a small-to-moderate number of knots. Among the quantile-score leaders, AR(2) specifications appear frequently, particularly at  $q(0.95)$  and above. MRTS specifications are more selective: they improve certain tail exceedance scores but are not uniformly preferred across all score types or quantile levels.

| Level  | Rank 1               | Rank 2               | Rank 3               | Rank 4               |
|--------|----------------------|----------------------|----------------------|----------------------|
| 0.9000 | 204 (MRTS, 0.041)    | 55 (Morris's, 0.041) | 205 (MRTS, 0.041)    | 70 (Morris's, 0.041) |
| 0.9500 | 70 (Morris's, 0.020) | 41 (Morris's, 0.020) | 12 (Morris's, 0.020) | 39 (Morris's, 0.020) |
| 0.9800 | 204 (MRTS, 0.008)    | 117 (AR2, 0.008)     | 67 (Morris's, 0.008) | 111 (AR2, 0.008)     |
| 0.9900 | 117 (AR2, 0.003)     | 15 (Morris's, 0.003) | 70 (Morris's, 0.003) | 61 (Morris's, 0.003) |
| 0.9950 | 15 (Morris's, 0.001) | 70 (Morris's, 0.001) | 61 (Morris's, 0.001) | 117 (AR2, 0.001)     |

Table 5: Validation selection result using 300 sites to predict 100 sites

| Level  | Rank 1               | Rank 2               | Rank 3               | Rank 4               |
|--------|----------------------|----------------------|----------------------|----------------------|
| 0.9000 | 54 (Morris's, 0.050) | 105 (AR2, 0.050)     | 35 (Morris's, 0.050) | 215 (MRTS, 0.050)    |
| 0.9500 | 117 (AR2, 0.030)     | 116 (AR2, 0.030)     | 114 (AR2, 0.030)     | 70 (Morris's, 0.030) |
| 0.9800 | 215 (MRTS, 0.015)    | 205 (MRTS, 0.015)    | 111 (AR2, 0.015)     | 58 (Morris's, 0.015) |
| 0.9900 | 58 (Morris's, 0.008) | 112 (AR2, 0.008)     | 216 (AR2, 0.008)     | 67 (Morris's, 0.008) |
| 0.9950 | 29 (Morris's, 0.004) | 35 (Morris's, 0.004) | 60 (Morris's, 0.004) | 116 (AR2, 0.004)     |

Table 6: Validation selection result using 200 sites to predict 200 sites

To gauge how stable these rankings are to the train/validation split, we repeat the comparison at two smaller training sizes. In the primary 400/400 split the empirical Brier thresholds at 0.90, 0.95, 0.98, 0.99, and 0.995 are 73.25, 79.75, 88.50, 94.75, and 100.52 ppb, and the top settings at these levels are numbered 204, 111, 58, 68, and 124 in the candidate list. When the split is changed to 300 training and 100 validation sites, the corresponding leaders become 204, 70, 204, 117, and 15; with 200 training and 200 validation sites they become 54, 117, 215, 58, and 29. The same broad model families (Morris skew- $t$ , AR(2), and MRTS variants) remain competitive across all three splits, but the exact best-ranked setting shifts with the split and the exceedance level. Such instability is expected: at the 0.99 and 0.995 levels only a handful of validation observations exceed the threshold, so the ranking is sensitive to which sites happen to fall in the held-out set. The real-data analysis therefore should not be read as identifying a single universally optimal specification, but as evidence that AR(2) and MRTS variants are valuable additions to the candidate pool when the inferential target is

held-out tail prediction.

Overall, the real-data analysis provides a stronger practical motivation for the proposed extensions than the simulation study alone. AR(2) yields additional temporal flexibility and appears among the leading models for several upper-tail quantile scores; MRTS enriches the spatial regression component and improves selected exceedance scores. The final modelling recommendation, accordingly, is not to replace the Morris baseline by a single proposed specification, but to include AR(2) and MRTS variants in the candidate pool and to select among them using held-out Brier and quantile scores on the application of interest.

## 5 MERRA-2 data analysis

To complement the ozone case study with a controlled medium-size benchmark, we conduct a preliminary analysis on hourly surface-temperature output from the MERRA-2 reanalysis. The objective is to assess whether the AR(2) temporal extension yields measurable predictive gains in a regime where the data are smoother and less tail-dominated than ground-level ozone. The data comprise  $n_s = 100$  spatial locations observed at  $n_t = 100$  consecutive hourly time points, with longitude and latitude retained for each site.

### 5.1 Data preprocessing and model setup

We randomly split the sites into training and validation subsets using seed 123, allocating 70% of the sites to model fitting and the remaining 30% to held-out evaluation. The regression design uses an intercept and a first-order spatial trend,

$$\mathbf{X}_t(\mathbf{s}) = [1, s_1, s_2]^\top, \quad (23)$$

where  $s_1$  and  $s_2$  are the longitude and latitude of site  $\mathbf{s}$ . Because the covariates are not standardized in this experiment, we set `nu.upper = 10` so that the smoothness parameter  $\nu$  of the Matérn covariance cannot drift into the numerically unstable large- $\nu$  regime during MCMC.

To isolate the effect of temporal structure, the three temporal specifications – no time series, AR(1), and AR(2) – share the same core configuration: a skew- $t$  margin, a Matérn covariance,  $K = 5$  knots, and no censoring threshold. Each chain is run for 20 000 iterations with the first 10 000 discarded as burn-in, yielding 10 000 retained posterior draws per model.

## 5.2 Results

Table 7 reports quantile-score (QS) and Brier-score (BS) comparisons across the three temporal specifications. The differences are small in absolute terms but qualitatively informative: the no-time-series model performs marginally better on the QS metrics and attains the lowest BS at the most extreme exceedance level (0.99), whereas AR(2) attains the best BS at the 0.90 and 0.95 levels. None of the three specifications therefore dominates uniformly; the preferred model depends on whether one targets bulk predictive accuracy, mid-tail exceedance probabilities, or the far tail.

The remaining summary metrics reinforce this reading. MAPE is slightly lower for AR(2) (0.0106), with AR(1) close behind, while the no-time-series model is somewhat worse on this bulk-accuracy measure. The 95% predictive-interval coverage is nearly identical across the three models (0.912–0.916), suggesting that temporal dependence affects predictive sharpness more than calibration. These modest gains come at a non-trivial computational cost: total fitting time rises from 777.3 seconds for the no-time-series model to 1041.7 seconds for AR(1) and 1075.8 seconds for AR(2). Taken together, the MERRA-2 results suggest that AR(2) is a defensible option when slightly better mid-tail behavior is desired, but the improvement is incremental rather than decisive – consistent with the picture obtained from the simulation study.

| Model       | QS90         | QS95         | QS99         | BS90         | BS95         | BS99         |
|-------------|--------------|--------------|--------------|--------------|--------------|--------------|
| No Temporal | <b>4.714</b> | <b>2.522</b> | <b>0.570</b> | 0.041        | 0.018        | <b>0.008</b> |
| AR(1)       | 4.716        | 2.522        | 0.570        | 0.040        | 0.016        | 0.009        |
| AR(2)       | 4.717        | 2.522        | 0.571        | <b>0.040</b> | <b>0.016</b> | 0.009        |

Table 7: Comparison of temporal specifications in the MERRA-2 experiment. Smaller values indicate better predictive performance.

## 6 Discussion

## References

Gneiting, Tilmann and Adrian E Raftery (Mar. 2007). “Strictly Proper Scoring Rules, Prediction, and Estimation”. In: *Journal of the American Statistical Association* 102.477.   
\_eprint: <https://doi.org/10.1198/016214506000001437>, pp. 359–378. ISSN: 0162-1459. DOI:

10.1198/016214506000001437. URL: <https://doi.org/10.1198/016214506000001437> (visited on 11/10/2025).

Morris, Samuel A. et al. (Sept. 2017). “A Space-Time Skew-t Model for Threshold Exceedances”.

In: *Biometrics* 73.3, pp. 749–758. ISSN: 0006-341X. DOI: 10.1111/biom.12644. URL: <https://doi.org/10.1111/biom.12644> (visited on 03/14/2025).

Tzeng, ShengLi and Hsin-Cheng Huang (Apr. 2018). “Resolution Adaptive Fixed Rank Kriging”. In: *Technometrics* 60.2. \_eprint: <https://doi.org/10.1080/00401706.2017.1345701>, pp. 198–208. ISSN: 0040-1706. DOI: 10.1080/00401706.2017.1345701. URL: <https://doi.org/10.1080/00401706.2017.1345701> (visited on 03/22/2026).

