

Non-stationary simulation study: what a mean-only MRTS augmentation can and cannot buy

July 2026 — interim report v2 (settings 1–12 and the threshold-vs-symmetry
disambiguation fitted)

Abstract

We evaluate the MRTS covariate extension of the Morris et al. (2017) skew- t space-time model on eleven simulated data-generating processes (DGPs) that place non-stationarity in different moments: the mean (settings 1, 2, 11), the spatial dependence (settings 3–8), the tails (setting 9), and the skewness (setting 10). Three findings. (i) MRTS improves predictive scores *only* where the truth carries time-averaged mean structure, and this is provable in advance: an oracle calculation shows the relative Brier score has a floor $R^* \approx 1$ for nine of eleven settings, so their observed nulls are properties of the design, not evidence about MRTS. (ii) Where mean structure exists, the energy- and variogram-score gains are carried *entirely* by the two thresholded symmetric- t methods. A disambiguating pair of methods (symmetric- t *without* threshold, $n = 10$) attributes this cleanly to the *threshold*, not the symmetry — and the absolute scores reframe it: POT censoring without a mean basis damages the predictive field so badly (baseline energy $2.6 \times$ – $16 \times$ worse) that MRTS is a partial *rescue*, not a synergy; the unthresholded methods never needed rescuing and remain better in absolute terms at every K . (iii) The Brier score is close to blind to mean recovery and can even deteriorate with K : in the strong-mean setting 11 a fit that recovers every parameter of the DGP at $K = 25$ has a *worse* Brier score than the badly misspecified $K = 0$ baseline. The rich-mean setting 12, run out to $K = 50$, resolves this: the deterioration is confined to intermediate K (where the mean basis is still too coarse), and once the basis is fine enough ($K \geq 30$ – 40) the Brier score improves by 7–12%, consistent with the extrapolation-error mechanism.

1 Design

Data-generating process. All settings share the skew- t -1 core on $n_s = 144$ sites drawn uniformly on $[0, 10]^2$:

$$Y_t(\mathbf{s}) = 10 + \lambda \sigma_t |z_t| + \sigma_t v_t(\mathbf{s}) + [\text{structure}]_k, \quad t = 1, \dots, 50, \quad (1)$$

with $\lambda = 3$, $z_t \stackrel{iid}{\sim} N(0, 1)$ half-normal, $v_t \sim N(0, C)$ exponential correlation ($\rho = 1$, $\gamma = 0.9$), $\tau_t \sim \text{Ga}(3/2, 4)$, $\sigma_t = \tau_t^{-1/2}$ (marginal t_3), i.i.d. in time. The settings differ only in the structure term—or, for the *replace- C* settings, in the error correlation itself (Table 1). Table 2 gives the resulting anatomy of each DGP: the marginal law of $Y_t(\mathbf{s})$, the spatial dependence of the error process, and the time-averaged true mean surface $\bar{\mu}(\mathbf{s})$ (the target of any pooled- β mean basis).

Fitted methods and grid. Five baseline methods are fitted to every dataset: (1) Gaussian; (2) skew- t , $K_{\text{knots}} = 1$; (3) symmetric- t , $K_{\text{knots}} = 1$, POT threshold at $q(0.80)$; (4) skew- t , $K_{\text{knots}} = 5$;

Table 1: The eleven fitted DGP settings and the number of simulation replicates n (independent datasets) fitted per setting. Routes: *mean* = added to the fixed mean; *additive* = mean-zero random field added to the draw; *replace C* = non-stationary correlation handed to the generator (unit diagonal, so the *marginal* skew- t law is unchanged — verified by KS tests, $p \in [0.075, 0.864]$); *skew term* = spatially varying λ .

ID	surf_type	non-stationarity	route	n
1	invariant	mean, 2 cosine bumps, sd = 3.7	mean	10
2	varying	mean, time-varying weights	mean	10^a
3	ns.dependence	dependence, low-rank cosine RE	additive	10
4	ps_cov	dependence, Paciorek–Schervish	replace C	3
5	ps_add	dependence, same PS field	additive	3
6	covdep_cov	dependence, $\rho(\mathbf{s}) = e^{0.7f_1(\mathbf{s})}$	replace C	3
7	covdep_add	dependence, same field	additive	3
8	fuentes_cov	dependence, 4-regime mixture	replace C	3
9	gh_marginal	tails, Tukey g -and- h , $g(\mathbf{s}), h(\mathbf{s})$	additive	3
10	lambda_varying	skewness, $\lambda(\mathbf{s}) = 3 + 3f_2(\mathbf{s})$	skew term	3
11	invariant_strong	mean, $3\times$ setting 1, sd = 11	mean	3

^a setting 2 was fitted at $K = 0$ only (no MRTS sweep); it is excluded from the score tables.

(5) symmetric- t , $K_{\text{knots}} = 5$, threshold $q(0.80)$. Each is run across MRTS design widths $K_{\text{MRTS}} \in \{0, 3, 5, 8, 10, 12, 15, 20, 25\}$, where $K_{\text{MRTS}} = 0$ denotes the baseline design $[1, s_1, s_2]$ and $K_{\text{MRTS}} = K$ replaces the design with the K -column `autoFRK::mrts` basis built on the *training* sites only (column 1 = intercept, 2–3 linear, 4: K thin-plate-spline eigenfunctions; $K = 3$ spans the baseline). MCMC: 20,000 iterations, 10,000 burn-in. The random seed is shared across K_{MRTS} within each (method, dataset), so all K -comparisons are *paired*.

Scoring. The last 44 of the 144 sites are held out entirely; all scores are computed out-of-sample on those sites. With $\hat{F}$ the posterior predictive law:

$$\text{BS} = \frac{1}{11} \sum_p \frac{1}{44 \cdot 50} \sum_{i,t} \left(\hat{P}(Y_t(\mathbf{s}_i^*) > u_p) - \mathbb{1}\{y_t(\mathbf{s}_i^*) > u_p\} \right)^2, \quad p \in \{0.90, \dots, 0.995\}, \quad (2)$$

$$\text{ES} = \frac{1}{50} \sum_t \left[\mathbb{E} \|\mathbf{X}_t - \mathbf{y}_t\| - \frac{1}{2} \mathbb{E} \|\mathbf{X}_t - \mathbf{X}'_t\| \right], \quad \mathbf{X}_t, \mathbf{X}'_t \sim \hat{F}_t \text{ (44-dim)}, \quad (3)$$

$$\text{VS} = \frac{1}{50} \sum_t \sum_{i < j} \left(|y_{it} - y_{jt}|^{1/2} - \mathbb{E} |X_{it} - X_{jt}|^{1/2} \right)^2. \quad (4)$$

The Brier score is a per-cell *marginal* score and cannot see spatial dependence; the energy score is the multivariate CRPS (reduces to CRPS at $d = 1$) and is location-sensitive; the variogram score ($p = 1/2$) targets pairwise dependence. Every result below is the *relative* score $R_{m,K,d} = S_{m,K,d}/S_{m,0,d}$ against the same method’s own $K = 0$ fit, averaged over datasets; $R < 1$ means MRTS helps.

2 Results

Tables 3–5 report, per setting and method, the best (smallest) $\bar{R}_{m,K}$ over $K > 0$, with the minimising K in parentheses. Asterisks mark $|\bar{R} - 1| > 2\widehat{\text{SE}}$ (SE over datasets; with $n = 3$ this is

Table 2: Anatomy of the twelve DGPs: marginal law of $Y_t(\mathbf{s})$, spatial dependence of the error process, and time-averaged true mean $\bar{\mu}(\mathbf{s})$ (its spatial sd in parentheses; per-dataset values for random means). $C_{\text{exp}}(h) = 0.9 e^{-h}$ is the shared stationary exponential correlation (Matérn $\nu = \frac{1}{2}$, $\rho = 1$, unit diagonal, i.e. a 10% nugget); “skew- t ” is the t_3 core of (1) with $\lambda = 3$; $*$ denotes convolution with an independent additive field ($\sigma_{\text{add}} = 3$); f_1, f_2 are the two cosine bumps.

ID	surf_type	marginal law	spatial dependence	mean structure $\bar{\mu}(\mathbf{s})$
1	invariant	skew- t	C_{exp} , stationary	$10 + 5f_1 + 3f_2$ (3.7)
2	varying	skew- t	C_{exp} , stationary	$10 + w_{1t}f_1 + w_{2t}f_2$, (w_{1t}, w_{2t}) $\sim$ $N(0, \text{diag}(25, 9))$ per day; $\bar{\mu} \approx 10$ (0.4)
3	ns_dependence	skew- $t * \text{Gaussian}$	$C_{\text{exp}} + \text{low-rank cosine}$ $\sum_j \tau_j^2 \phi_j(\mathbf{s}) \phi_j(\mathbf{s}')$; non- stationary, negative at long range	10 (RE mean-zero) (0.4)
4	ps_cov	skew- t (unchanged)	Paciorek–Schervish locally anisotropic kernel; scale, as- pect and rotation each vary across the domain	10 exactly (0)
5	ps_add	skew- $t * \text{Gaussian}$	$C_{\text{exp}} + \sigma_{\text{add}}^2 C_{\text{PS}}$, additive	10 (RE mean-zero) (0.4)
6	covdep_cov	skew- t (unchanged)	isotropic PS kernel with lo- cal range $\rho(\mathbf{s}) = e^{0.7f_1(\mathbf{s})} \in$ $[0.5, 2.0]$	10 exactly (0)
7	covdep_add	skew- $t * \text{Gaussian}$	$C_{\text{exp}} + \sigma_{\text{add}}^2 C_{\rho(\mathbf{s})}$, additive	10 (RE mean-zero) (0.4)
8	fuentes_cov	skew- t (unchanged)	Fuentes 4-regime weighted GP mixture, ranges 0.4–4	10 exactly (0)
9	gh_marginal	skew- $t * \text{std. Tukey } g\text{-}$ $\text{and-}h$; $g(\mathbf{s}) = 0.8f_1$, $h(\mathbf{s}) = 0.12(1+f_2)$	C_{exp} in both components (transform perturbs correla- tion only mildly)	10 (standardised mean- zero) (0.4)
10	lambda_varying	skew- t with $\lambda(\mathbf{s}) = 3 +$ $3f_2 \in [0, 6]$	C_{exp} , stationary	induced: $10 + \lambda(\mathbf{s}) \bar{z}$, $\bar{z} =$ $\frac{1}{T} \sum_t \sigma_t z_t $ (2.8)
11	invariant_strong	skew- t	C_{exp} , stationary	$10 + 15f_1 + 9f_2$ (11.0)
12	mrts_multibump	skew- t	C_{exp} , stationary	10 + six alternating-sign Gaussian bumps (width 1.6) (11.0)

indicative only). Table 6 gives the full setting $\times K$ profile of the relative Brier score, method by method.

2.1 Finding 1: gains appear only where the truth has mean structure — and this is provable a priori

The only large entries in any table are settings 1 and 11 (mean structure), methods 3 and 5: energy and variogram reductions of 34–42% (setting 1) and 73–80% (setting 11). Every dependence, tails and skewness setting is flat to within a few percent, in *all three* scores.

This dichotomy can be established without MCMC. Because the MRTS coefficient β is pooled over time, the most any mean basis can deliver is the time-averaged true mean $\bar{\mu}(\mathbf{s})$. Define the oracle relative Brier

$$R^* = \frac{\text{BS}(\bar{\mu})}{\text{BS}(\text{spatial constant})}, \quad (5)$$

both evaluated under the *true* skew- t predictive law. R^* is a floor: no K , sample size or sampler can

Table 3: **Energy score**, best over $K > 0$ relative to the same method’s $K = 0$ (minimising K in parentheses). Values < 1 : MRTS helps.

ID	non-stationarity	n	Gaussian	Skew- t K1	Sym- t K1 T	Skew- t K5	Sym- t K5 T
1	mean (2 bumps, sd 3.7)	10	0.995* (25)	0.982* (25)	0.579* (20)	0.987* (20)	0.659* (12)
3	dependence (low-rank)	10	1.000 (5)	1.000 (3)	0.993 (5)	1.000 (15)	0.911 (3)
4	dependence (PS, repl. C)	3	0.999* (5)	1.000 (8)	0.943* (8)	0.996* (3)	0.846 (15)
5	dependence (PS, additive)	3	0.999 (12)	0.999* (8)	0.953 (15)	0.998* (10)	0.955 (15)
6	dependence ($\rho(\mathbf{s})$, repl. C)	3	1.000 (10)	0.999 (10)	0.861* (20)	0.997* (12)	0.941* (10)
7	dependence ($\rho(\mathbf{s})$, additive)	3	1.000 (25)	1.000 (10)	0.988 (20)	0.995 (25)	0.926 (15)
8	dependence (Fuentes, repl. C)	3	1.000 (10)	1.000 (3)	0.848 (25)	0.998 (5)	0.721* (10)
9	tails (g -and- h)	3	0.999* (5)	1.000 (10)	0.887 (5)	0.986 (25)	0.942 (5)
10	skewness ($\lambda(\mathbf{s})$)	3	0.994 (25)	1.000 (3)	0.975 (5)	0.998 (25)	0.938* (5)
11	mean (2 bumps, sd 11)	3	0.990 (25)	0.984 (25)	0.199* (20)	0.982 (25)	0.267* (20)

Table 4: **Brier score** (mean over the 11 exceedance levels), best over $K > 0$ relative to own $K = 0$.

ID	non-stationarity	n	Gaussian	Skew- t K1	Sym- t K1 T	Skew- t K5	Sym- t K5 T
1	mean (2 bumps, sd 3.7)	10	1.000 (3)	0.985* (25)	0.994* (25)	0.993 (20)	0.971 (3)
3	dependence (low-rank)	10	1.000 (5)	0.999 (5)	1.000 (3)	0.997 (15)	0.993 (12)
4	dependence (PS, repl. C)	3	0.999 (25)	0.999 (25)	0.998* (5)	0.985* (25)	0.940* (20)
5	dependence (PS, additive)	3	1.000 (3)	0.999 (12)	1.000 (8)	0.990 (15)	0.973 (5)
6	dependence ($\rho(\mathbf{s})$, repl. C)	3	0.999 (20)	0.999* (5)	0.999* (15)	0.983* (12)	0.970 (10)
7	dependence ($\rho(\mathbf{s})$, additive)	3	1.000 (15)	0.999 (10)	0.998 (15)	0.990 (5)	0.968 (15)
8	dependence (Fuentes, repl. C)	3	1.000* (3)	1.000 (3)	0.999 (3)	0.997 (12)	0.962 (12)
9	tails (g -and- h)	3	1.000 (12)	1.000 (15)	1.000 (3)	0.996 (10)	0.967 (12)
10	skewness ($\lambda(\mathbf{s})$)	3	0.995 (25)	0.999 (3)	0.957* (25)	0.994 (25)	0.988 (20)
11	mean (2 bumps, sd 11)	3	1.000 (3)	1.000 (3)	0.959* (25)	1.004 (3)	0.997 (3)

push the relative Brier below it. Table 7 shows $R^* \approx 1$ for nine of eleven settings: their Brier nulls are design properties, not evidence about MRTS. Setting 10 is the striking case: $R^* = 1.12 > 1$ — supplying the *true* mean makes the Brier score *worse*, because the mean shift induced by $\lambda(\mathbf{s})$ is a symptom of a skewness change that the constant- λ model cannot represent; correcting the symptom while the disease remains actively misleads the tail probabilities.

2.2 Finding 2: the gains are carried entirely by the thresholded methods

In Table 3, setting 11 (mean sd = 11), the two POT-thresholded symmetric- t methods reduce the energy score by 73–80%, while the Gaussian and both skew- t methods stay within 2% of their $K = 0$ baselines — despite those same skew- t fits recovering the mean surface almost exactly (recovery RMSE $14.0 \rightarrow 0.66$ for method 2 at $K = 25$) and recovering every DGP parameter (Section 2.3). Table 8 shows the full K profile.

Methods 3/5 are simultaneously (a) POT-thresholded and (b) symmetric- t , so a disambiguating pair was added: methods 11/12, *identical to methods 3/5 except threshold = 0*, fitted on setting 1 with $n = 10$ (Table 9). The verdict is unambiguous: the unthresholded symmetric- t methods are *flat* (energy ratio 0.99 at every K , indistinguishable from the skew- t reference), while their thresholded twins gain 35–40%. **The threshold, not the symmetry, is the driver.**

The absolute column reframes the finding. The thresholded baselines are *catastrophically* worse to begin with: censoring everything below $q(0.80)$ leaves the constant-mean fit with almost no

Table 5: **Variogram score** ($p = 1/2$), best over $K > 0$ relative to own $K = 0$.

ID	non-stationarity	n	Gaussian	Skew- t K1	Sym- t K1 T	Skew- t K5	Sym- t K5 T
1	mean (2 bumps, sd 3.7)	10	0.995* ⁽²⁵⁾	0.976* ⁽²⁵⁾	0.617* ⁽¹⁵⁾	0.984* ⁽²⁰⁾	0.712 ⁽¹²⁾
3	dependence (low-rank)	10	1.000 ⁽⁵⁾	1.000 ⁽³⁾	0.984 ⁽⁵⁾	0.999 ⁽²⁵⁾	0.894 ⁽³⁾
4	dependence (PS, repl. C)	3	0.999* ⁽¹⁰⁾	1.000 ⁽¹⁵⁾	0.988 ⁽¹⁵⁾	0.997 ⁽¹⁰⁾	0.990 ⁽³⁾
5	dependence (PS, additive)	3	0.999* ⁽¹²⁾	0.999* ⁽⁸⁾	0.971 ⁽¹⁵⁾	0.998 ⁽¹⁵⁾	0.986 ⁽¹⁵⁾
6	dependence ($\rho(\mathbf{s})$, repl. C)	3	0.999 ⁽¹⁰⁾	0.999 ⁽¹⁰⁾	0.924* ⁽²⁰⁾	0.997 ⁽¹²⁾	0.984 ⁽¹⁰⁾
7	dependence ($\rho(\mathbf{s})$, additive)	3	1.000 ⁽¹⁰⁾	1.000 ⁽⁸⁾	0.995 ⁽³⁾	0.995 ⁽⁵⁾	0.960 ⁽¹⁵⁾
8	dependence (Fuentes, repl. C)	3	1.000 ⁽⁸⁾	1.000 ⁽³⁾	0.890 ⁽²⁵⁾	1.000 ⁽³⁾	0.847 ⁽⁵⁾
9	tails (g -and- h)	3	1.000 ⁽³⁾	1.000 ⁽⁵⁾	0.865 ⁽⁵⁾	0.982 ⁽²⁵⁾	0.947 ⁽⁵⁾
10	skewness ($\lambda(\mathbf{s})$)	3	0.991 ⁽²⁵⁾	1.000 ⁽²⁵⁾	1.000 ⁽³⁾	0.995 ⁽²⁵⁾	0.905* ⁽¹⁰⁾
11	mean (2 bumps, sd 11)	3	0.973 ⁽²⁵⁾	0.953* ⁽²⁵⁾	0.201* ⁽¹⁰⁾	0.952* ⁽²⁵⁾	0.244* ⁽²⁰⁾

information about the spatial level of the field, and its out-of-sample predictive field degrades by a factor 2.6 (K1) to 16 (K5) in energy score. MRTS repairs the mean that the censored likelihood cannot learn — a partial *rescue*, not a synergy: even at the best K , method 3 reaches $21.8 \times 0.579 \approx 12.6$, still well above the 8.5 that methods 2/11 achieve *without any MRTS at all*. In absolute terms the unthresholded methods win at every K ; the correct reading of Finding 2 is therefore “*only the thresholded methods have room to be helped*,” and the practical implication is for POT-style threshold models specifically: if the mean is not modelled well, censored fitting amplifies the damage, and a mean basis buys back part of it.

2.3 Finding 3: the Brier score is nearly blind to mean recovery, and can punish it

Setting 11, method 2 (skew- t K1), dataset 1, posterior means (truth: $\lambda = 3$, $\mathbb{E}[\sigma] = 2.26$, $\rho = 1$, $\gamma = 0.9$):

K_{MRTS}	recovery RMSE	$\hat{\lambda}$	$\mathbb{E}[\hat{\sigma}]$	$\hat{\rho}$	$\hat{\gamma}$
0	14.01	0.15	11.16	2.04	0.99
10	2.51	2.04	3.01	1.21	0.91
25	0.66	3.18	1.94	0.96	0.94

At $K = 0$ the model cannot represent the sd-11 mean surface and absorbs the spatial variance into an inflated σ ($5 \times$ truth) with λ crushed to zero. At $K = 25$ *every* parameter is recovered. Yet the $K = 25$ Brier score is *worse* than $K = 0$ by 8.6% (Table 8), and the same reversal holds for methods 1, 4, 5. A correctly specified model loses to a badly misspecified one on this score. The working hypothesis — not yet verified — is an out-of-sample interaction: the sharpened predictive distribution ($\hat{\sigma}$ correct instead of inflated) exposes the residual error of the MRTS mean *extrapolated* to the held-out sites, whereas the misspecified $K = 0$ fit hides that error inside an accidentally conservative (inflated- σ) predictive distribution.

Setting 12 resolves the anomaly. The rich-mean setting 12 (six bumps, sd 11) was run out to $K_{\text{MRTS}} = 50$, where the noise-free MRTS projection error of its surface finally becomes small (0.71 at $K = 5$, 0.13 at $K = 25$, 0.057 at $K = 50$). The Brier profile follows exactly the pattern the extrapolation mechanism predicts: flat-to-worse at intermediate K (ratios 1.00–1.05 for $K \leq 25$, where the basis is still too coarse and the residual mean error is exposed by the sharpened predictive), then *improving* once the basis is fine enough:

Table 6: **Relative Brier score at every K_{MRTS}** (vs the same method’s own $K = 0$, mean over datasets; * marks $|\bar{R} - 1| > 2\widehat{\text{SE}}$). Setting 2 is excluded ($K = 0$ only); $K \geq 30$ was fitted for setting 12 only. Isolated large values of Sym- t K5 T are the exponential-fallback outliers discussed in the Caveats and are reported as-is.

ID	n	$K=3$	$K=5$	$K=8$	$K=10$	$K=12$	$K=15$	$K=20$	$K=25$	$K=30$	$K=40$	$K=50$
<i>Gaussian</i>												
1	10	1.000	1.016*	1.023*	1.022*	1.019*	1.016*	1.009*	1.008*	—	—	—
3	10	1.000	1.000	1.001*	1.001*	1.000	1.000	1.000	1.001	—	—	—
4	3	1.000	0.999*	0.999*	1.000	1.000	1.000	0.999	0.999	—	—	—
5	3	1.000	1.001	1.001	1.000	1.001*	1.000	1.001*	1.002*	—	—	—
6	3	1.000	1.000	1.000	1.000	1.000	1.000	0.999	1.000	—	—	—
7	3	1.000	1.000	1.000	1.000	1.000	1.000	1.002*	1.001	—	—	—
8	3	1.000*	1.001	1.001	1.001	1.000	1.000	1.002	1.002	—	—	—
9	3	1.000*	1.000	1.000	1.000*	1.000	1.000	1.000	1.001	—	—	—
10	3	1.000	1.003	1.008	1.004	1.000	0.997	0.995	0.995	—	—	—
11	3	1.000	1.059*	1.078*	1.101*	1.114*	1.095*	1.069*	1.051*	—	—	—
12	3	1.000	1.003*	1.000	0.993	0.980	1.019	1.049	1.045	0.988	0.947	0.930
<i>Skew-t K1</i>												
1	10	1.000	0.996	1.006	1.006	1.001	0.996	0.987*	0.985*	—	—	—
3	10	1.000	0.999	1.000	1.000	1.000	1.000	1.000	1.000	—	—	—
4	3	1.000*	1.000	0.999	1.000	1.000	1.000	0.999	0.999	—	—	—
5	3	1.000	1.000	0.999	1.000	0.999	1.000	1.000	1.001	—	—	—
6	3	1.000	0.999*	0.999*	1.000	1.001*	0.999	1.000	1.000	—	—	—
7	3	1.000	1.000	1.001	0.999	1.000	1.000	1.001	1.000	—	—	—
8	3	1.000	1.000	1.001	1.000	1.001	1.001*	1.001	1.000	—	—	—
9	3	1.001	1.001	1.001*	1.001	1.001	1.000	1.001	1.001	—	—	—
10	3	0.999	1.003	1.008*	1.008	1.005	1.003	1.002	1.001	—	—	—
11	3	1.000	1.052*	1.074*	1.103*	1.163*	1.138*	1.110*	1.086*	—	—	—
12	3	1.000	1.005*	0.999	0.992	0.965	1.022	1.024	1.020	0.980	0.946	0.925*
<i>Sym-t K1 T</i>												
1	10	1.000	1.003	1.006	1.008*	1.003	1.001	0.996	0.994*	—	—	—
3	10	1.000	1.000	1.001	1.001	1.001	1.001	1.001	1.001	—	—	—
4	3	1.000	0.998*	0.999	0.999	1.000	0.999	1.000	0.999	—	—	—
5	3	1.000*	1.000	1.000	1.000	1.000	1.001	1.000	1.001	—	—	—
6	3	1.000	1.002*	1.000	1.000	1.001	0.999*	0.999	1.000	—	—	—
7	3	1.000*	1.000	1.000	1.000	1.000	0.998	1.001	1.000	—	—	—
8	3	0.999	0.999	0.999	1.002	1.000	1.001	1.003	1.008	—	—	—
9	3	1.000	1.001	1.001	1.000	1.001	1.001*	1.000	1.001*	—	—	—
10	3	0.999	0.993*	0.980*	0.985	0.975	0.970	0.963*	0.957*	—	—	—
11	3	0.999	1.016	1.021	1.039*	1.008	0.982	0.977*	0.959*	—	—	—
12	3	1.000	1.006*	0.997	0.907*	0.893*	0.899*	0.896*	0.918*	0.892*	0.896*	0.880*
<i>Skew-t K5</i>												
1	10	0.995	1.004	1.015*	1.008	1.010	1.011*	0.993	0.994	—	—	—
3	10	0.999	1.004	1.006	0.998	1.001	0.997	1.000	0.998	—	—	—
4	3	0.989*	0.993	0.988	0.994	0.992	0.994	0.987	0.985*	—	—	—
5	3	0.996	0.994	0.992*	0.992	0.994	0.990	0.993	0.997	—	—	—
6	3	0.998	1.001	1.000	0.996	0.983*	0.996	0.986*	0.989	—	—	—
7	3	0.993	0.990	0.997	0.994	0.994	0.994	0.994	0.993	—	—	—
8	3	0.999	0.999	1.007	1.011*	0.997	1.002	1.015*	1.007*	—	—	—
9	3	1.005	1.007	0.999	0.996	1.007	1.001	1.003	1.003	—	—	—
10	3	0.999	1.002*	1.006*	1.008	1.004	1.005	1.002	0.994	—	—	—
11	3	1.004	1.088*	1.096*	1.123*	1.184*	1.155*	1.142*	1.112	—	—	—
12	3	0.992	1.034	0.995	1.004	0.962	0.976	1.017	1.022	0.965	0.927*	0.913*
<i>Sym-t K5 T</i>												
1	10	0.971	1.057	1.073	1.097*	1.055	1.023	1.050	1.085	—	—	—
3	10	0.997	1.011	1.001	1.009	0.993	1.013	1.068	1.066	—	—	—
4	3	0.981*	0.980	0.999	0.998	1.022 6	0.994	0.940*	0.995	—	—	—
5	3	1.021	0.973	0.975	0.974	0.987	0.980	0.977	0.973	—	—	—
6	3	0.980	0.978	0.994	0.970	0.988	1.020	0.997	1.028	—	—	—
7	3	0.981	0.990	0.987	0.972	0.993	0.968	1.011	0.993	—	—	—
8	3	0.967	1.005	0.990	0.990	0.962	1.010	1.062*	1.050	—	—	—

Table 7: Oracle Brier floor R^* per setting (`oracle_brier_ceiling.R`; Monte Carlo, 3 datasets).

setting	$\text{sd}(\bar{\mu})$	R^*	verdict
1 invariant	3.68	0.856	weak headroom (14%)
2 varying	0.42	1.000	dead end
3 ns.dependence	0.41	0.999	dead end
4 ps.cov	0.00	1.000	dead end
5 ps.add	0.39	0.999	dead end
6 covdep.cov	0.00	1.000	dead end
7 covdep.add	0.41	1.000	dead end
8 fuentes.cov	0.00	1.000	dead end
9 gh.marginal	0.43	1.000	dead end
10 lambda.varying	2.80	1.120	harmful
11 invariant_strong	11.04	0.439	positive control
12 mrts_multibump	11.00	0.495	positive control

Table 8: Setting 11, relative score vs K_{MRTS} ($n = 3$): the thresholded symmetric- t methods against the unthresholded skew- t .

score	method	$K=3$	$K=5$	$K=8$	$K=10$	$K=12$	$K=15$	$K=20$	$K=25$
energy	Skew- t K1 (no thr.)	0.999	1.013	1.022	1.027	1.023	1.017	0.996	0.984
	Sym- t K1 (thr.)	0.990	0.263	0.266	0.229	0.201	0.244	0.199	0.256
	Sym- t K5 (thr.)	1.001	0.317	0.539	0.374	0.342	0.462	0.267	0.333
Brier	Skew- t K1 (no thr.)	1.000	1.052	1.074	1.103	1.163	1.138	1.110	1.086
	Sym- t K1 (thr.)	0.999	1.016	1.021	1.039	1.008	0.982	0.977	0.959
	Sym- t K5 (thr.)	0.997	1.078	1.182	1.114	1.163	1.120	1.104	1.133

relative Brier, setting 12	$K=15$	$K=25$	$K=30$	$K=50$	recovery at $K=50$
Gaussian	1.019	1.045	0.988	0.930	0.512
Skew- t K1	1.022	1.020	0.980	0.925	0.073
Sym- t K1 thr.	0.899	0.918	0.892	0.880	0.422
Skew- t K5	0.976	1.022	0.965	0.913	0.268

(Method 5 is excluded: its exponential-fallback instability blows up one of the three datasets at large K , ratio 2.7.) The $K=50$ gains of 7–12% are real but far above the oracle floor $R^* = 0.495$ — the mean channel is recovered (skew- t K1 recovery ratio 0.073), yet most of the theoretically available Brier headroom is spent compensating the sharpness–extrapolation trade-off. Practically: *for judging mean-structure recovery, the Brier score remains the wrong instrument — it responds only when the basis is fine enough to make the residual extrapolation error negligible; energy/variogram scores (or the recovery RMSE directly) are the informative ones.*

3 Caveats

- Settings 4–11 are pilots with $n = 3$; settings 1 and 3 have $n = 10$. Starred cells with $n = 3$ should be read as indicative. The small- n *best-over- K* statistic also has a selection bias toward values < 1 ; flat profiles (e.g. Sym- t K5 rows in dependence settings, best ≈ 0.72 – 0.85 with non-monotone K profiles) are consistent with noise plus this selection effect.

Table 9: Disambiguation on setting 1 ($n = 10$): relative energy score vs K_{MRTS} , and the *absolute* energy score at $K = 0$. Method pairs (3, 11) and (5, 12) differ *only* in the POT threshold.

	method	$K=5$	$K=10$	$K=15$	$K=20$	$K=25$	abs. ES at $K=0$
3	Sym- t K1, thr. $q(0.80)$	0.630	0.617	0.589	0.579	0.596	21.8
11	Sym- t K1, no threshold	0.996	0.997	0.991	0.990	0.989	8.5
5	Sym- t K5, thr. $q(0.80)$	0.688	1.113	0.773	0.718	0.697	137.1
12	Sym- t K5, no threshold	0.994	1.008	0.993	0.997	0.997	9.3
2	Skew- t K1 (reference)	0.994	0.994	0.985	0.983	0.982	8.5

- Method 5 (Sym- t K5 T) occasionally triggers its exponential-covariance fallback and produces isolated outlier fits (e.g. setting 4, $K = 25$: energy ratio 15.8 in one dataset). Best-over- K cells are robust to this; per- K curves at $n = 3$ are not.
- All K comparisons are paired (shared seeds across K), so within-method ratios are considerably more precise than the between-method comparisons.
- The recovery RMSE (not tabulated here) has a known bimodality on flat-mean settings caused by the intercept- λ ridge; see `nonsta_settings_design.md`.

4 The setting-12 showcase: why a large K is worth having

Setting 12’s six-bump surface was designed so that the MRTS approximation improves *gradually* with K (noise-free projection error 0.71 at $K=5$, 0.19 at $K=15$, 0.057 at $K=50$), in contrast to the 2-bump settings which the basis captures by $K=5$. The fitted results track that curve: skew- t K1 recovery ratios are 0.93 ($K=5$) $\rightarrow$ 0.49 ($K=10$) $\rightarrow$ 0.14 ($K=25$) $\rightarrow$ 0.073 ($K=50$) — the extended elbow, realised in MCMC. Two by-products: (a) the Brier anomaly resolution above; (b) a resolution ceiling — a higher-frequency mean component (3 cycles per axis) is *unrecoverable* on the 100-train/44-test geometry (its projection error *rises* with K), i.e. MRTS’s usable resolution is capped by site density, and the six-bump surface sits just inside that cap.

5 In progress

1. **Non- t extension (settings 13–21)**: all settings so far share the skew- t core, so at large K the fitted skew- t is *correctly specified* — a confound for every “MRTS helps/hurts” conclusion. The extension drops the skew- t generator entirely: nine nonlinear mean surfaces (Franke, curved ridge, annulus, sigmoid front, GP draw, polynomial, transform mixtures, non-smooth hinges; orthogonalised against the $K = 0$ design and standardised to sd 11) with Gaussian-GP errors, calibrated so that $R^* \approx 0.45$ throughout. Data generation in `setup_nonsta.R` $\rightarrow$ `simdata_nonsta_ext.RData`.
2. **n upgrade**: settings 4–12 are $n = 3$ pilots; the headline contrasts (11, 12, and the disambiguation) merit $n = 10$.